

Lecture 1



Computer Systems for Data Science

Topic 1

Course Introduction
Systems concepts



Agenda

- Intro to instructors
- High-level overview
 - What is data science and big data?
 - Class goals and why should you care?
- Class logistics
 - How the class is going to work?
- Performance and systems rules of thumb
- Intro to datacenters

Who Are We?



Course Instructors and TAs

- Instructor: Asaf Cidon
 - New Columbia EE + CS professor
 - Research focus: infrastructure for big data, storage systems for big data
 - Before Columbia, founded Sookasa, cloud storage security company (acquired by Barracuda Networks)
 - Built big data systems and used them both in academia and in industry
 - Type of problems I've worked on:
 - Automatic classification of spear phishing
 - Anomaly detection of users within organization
 - Sensitive data exfiltration...
- Head TA: Hongyi Wang
- TAs: Yu Jian Wu, Ke Li, Mingen Pan, Qianrui Zhang
- All TAs have experience in databases and systems
 - Hongyi and Yu Jian helped create course homework

What is Data Science and Big Data?



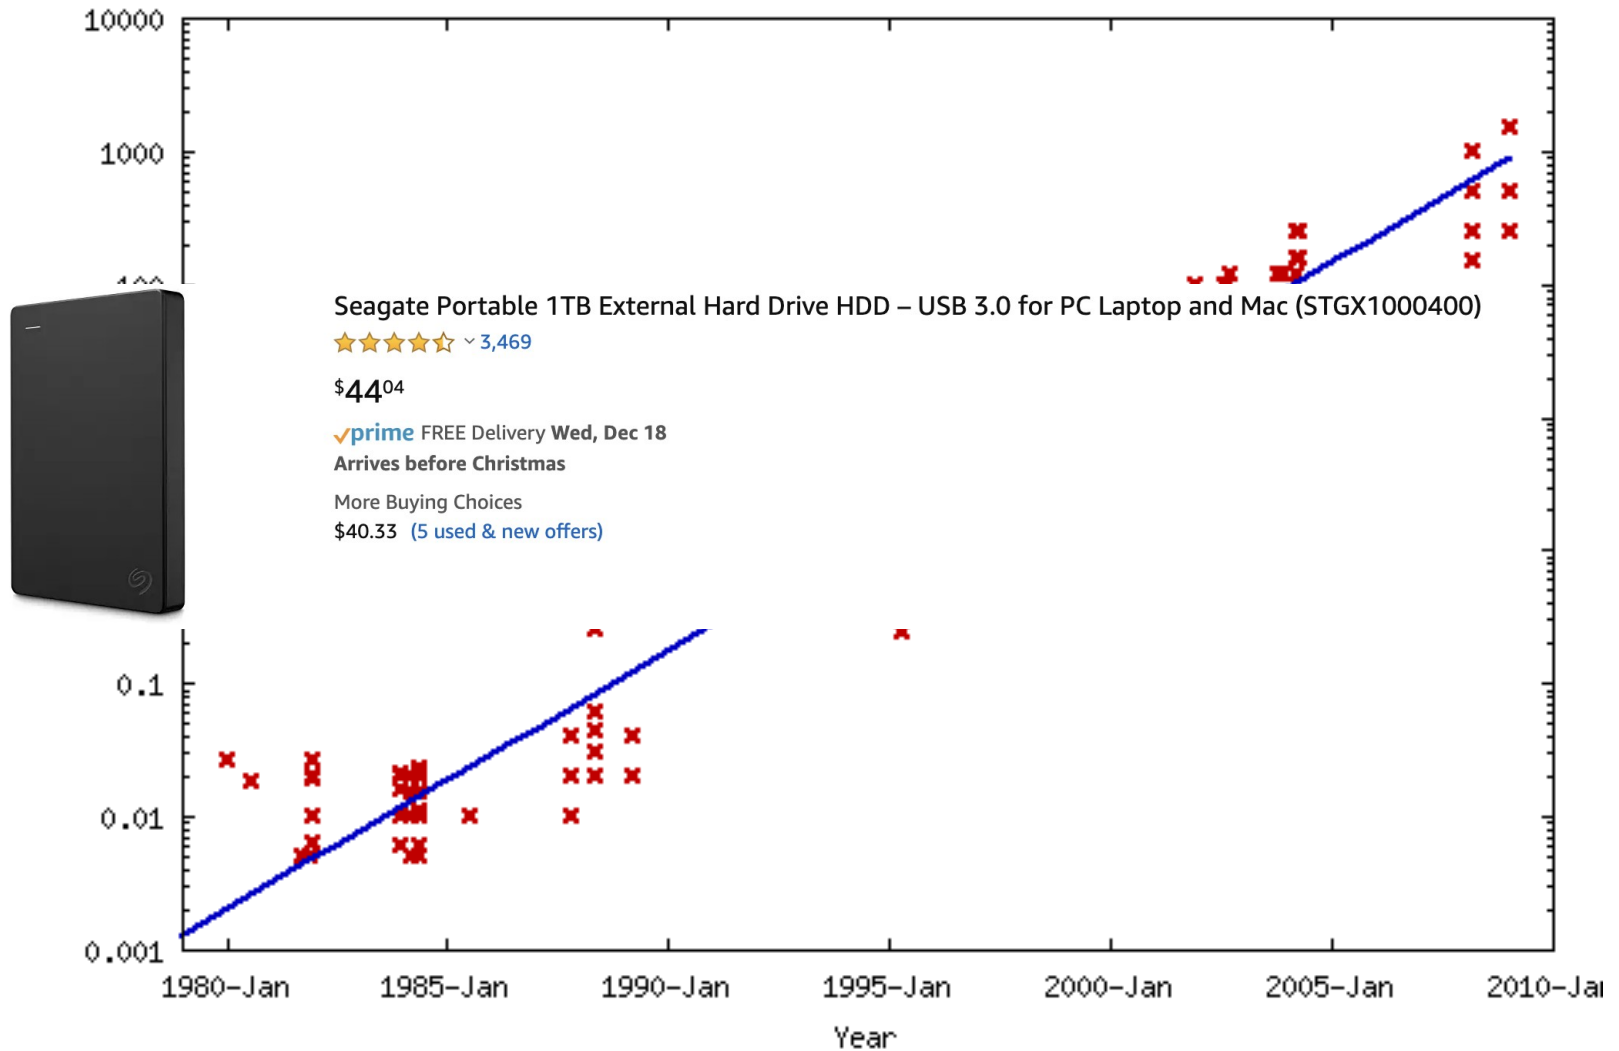
This was a system for big data

		67
June 11 ¹⁹²⁸	Geo. A. Kelly	Phoenix, Arizona.
June 14	Mrs. Chas. Wright	Phoenix Arizona
June 16	Nellora Wright	Phoenix Arizona
June 16	Charity A. Bales	Pescott - Arizona.
" "	Mrs. M. A. Carpenter	San Francisco
" "	Mr. & Mrs. Carpenter	Pescott
July 10	James Ostrom	251 S. Mt. Vernon St. Pescott
July 10	A. W. Gernung	Deerway Arizona
July 10	Millicent Gernung	Deerway, Arizona
" "	Walt Kulin	Ph. 1197 - N. Y.
July 11	Mrs. Raw. & Daughters	San Francisco Calif.
" "	Mrs. Ralph Roberts	Pescott
" "	Mrs. A. H. Favours	"
" "	Mrs. J. A. Miller	"
" "	Mrs. J. G. Harris	"
" "	Mrs. C. A. Hula	"
" "	Mary B. Soyars	"
" "	Mrs. & Baby Hoffman	"
" "	Mrs. (Kee) Young	"
" "	Mrs. A. L. Whitney	"
" "	Mrs. J. W. Stara	"
" "	Mrs. H. Robinson	"
		Arizona

Data science systems were expensive



Today: data is cheap



Where is data coming from?

- Physical devices



Where is data coming from?

- Physical devices
- Software logs

Where is data coming from?

- Physical devices
- Software logs
- Phones



Where is data coming from?

- Physical devices
- Software logs
- Phones
- GPS/Cars



Where is data coming from?

- Physical devices
- Software logs
- Phones
- GPS/Cars
- Internet of *Things*



Examples of data science

- What video should I recommend to this user to view next?
- Does this MRI image of a breast contain a tumor?
- Who is going to win the election?
- Which cities in the US will have high incidence of flu in 2 weeks?
- Is the object across from the car a pedestrian?

What is big data?

- “**Extremely large data sets** that may be analyzed computationally to reveal patterns, trends, and associations, especially relating to human behavior and interactions” – Oxford Dictionary
- What's an extremely large data set?
 - Fits on a single machine?
 - Fits on 10 machines?

Class Goals



Our focus in this class: **Computer Systems** for Data Science

- Questions we **will** answer in this class:

How are big data systems designed?

How to store the data?

How to query/analyze the data?

How do we ensure uptime/availability to the data?

How to ensure privacy/security/quality?

- Questions we **won't** answer in this class:

What algorithm should we use?

How to use ML for big data?

How do we explain/debug ML models?

How can data be visualized?

What are the statistical/mathematical foundations for data science?

Course Objectives

- **Graduate level course**
- **Broad overview of cloud systems that are used in data science**
 - **Database** related topics (DBMS, SQL, NoSQL)
 - **Computer systems** foundations (throughput vs. latency, scalability vs. performance)
 - **Distributed systems** for data scientists (sharding, fault tolerance)
- Throughout the class we will focus on how **commonly used and modern** cloud-based big data systems work (Spark, Tensorflow, Hadoop file system,...)
- The class will give a **broad and hopefully practical** introduction to these topics geared towards data scientists, but **does not replace** core CS/EE classes like OS, databases, distributed systems, security, architecture
- **You come from diverse backgrounds:** Some of the content will be repetitive for students who have taken the classes above
- **Required background**
 - Programming experience with Python or equivalent
 - Programming assignments 2-3 will be in pairs in Python, so we will make sure that all pairs have at least one student who knows Python

Course Administration and Grading

- **All materials, assignments, etc. posted on course website**
 - <https://cs-w4121.github.io/>
- **Announcement will be posted on Piazza, courseworks**
- **Lecture Materials**
 - Lecture slides
 - No textbook (new, fast moving field)
- **Homework, assignments, exams**
 - Programming assignment 1: BigQuery (20%), alone
 - Written assignment: systems and databases (10%), alone
 - Midterm (15%)
 - Programming assignment 2: Spark (20%), in pairs
 - Programming assignment 3: Tensorflow (10%), in pairs
 - Final exam (25%)

Programming Assignments

- 3 programming assignments
 - Assignment 1 is solo, assignments 2-3 are in pairs
- Programming assignments 2-3 are in Python
 - If you don't know Python
 - Pair with someone who does
 - A useful language to learn!
- All assignments in Google Cloud (GCP)
 - Goal: familiarize yourself with working in public cloud environment
 - AWS / Azure / GCP are similar
 - Many systems and deployment details are hidden / automated (but we won't ignore them!)
 - We will be focusing on systems-level problems, not on algorithms
 - We will provide GCP credits, if you run out contact us
 - If you reach \$10 of credits or less, please contact Mingen: pan.mingen@columbia.edu
 - But be careful not to spend too many!
- Programming assignment goals
 - Assignment 1: BigQuery
 - Learning to use SQL on a big data set
 - Assignment 2: DataProc + Spark + HDFS + Streaming
 - Understanding Spark, Streaming systems concepts
 - Assignment 3: Tensorflow
 - Understanding Tensorflow and basic ML training systems problems

Tentative Contents and Syllabus

- Computer systems and performance rules of thumb
 - Latency vs. throughput
 - Amdahl's law
 - Back-of-the-envelope systems math
 - Performance bottlenecks
- Data centers
 - What is a data center?
 - Data center failures
 - Achieving reliability with smart software
- Databases
 - Relational model and SQL
 - SELECT, FROM, WHERE
 - GROUPBY
 - JOINS
 - Nested queries
 - Transactions
 - ACID
 - OLAP vs. OLTP, SQL vs. NoSQL
 - Indexing
 - Logging
 - **System highlight: BigQuery, MySQL**
- Storage and distributed file systems
 - Storage technologies primer
 - Distributed file systems
 - **System highlight: Hadoop File System (HDFS), amazon S3/Google Cloud Storage**

Tentative Contents and Syllabus

- Distributed systems
 - 2 Phase Commit
 - Locking
 - Sharding
 - Fault tolerance
 - Replication and consensus
- Mapreduce
 - Mapreduce computing model
 - Stragglers
 - Importance of 99th latency
 - Strategies to mitigate tail latency and increase availability
- Distributed analytics and streaming
 - Resilient Distributed Dataframes (RDD)
 - Fault tolerance in distributed analytics: lineage
 - Streaming computing model
 - **Systems highlight: Spark, Google Dataproc, Spark streaming**
- Caching
 - Performance benefits
 - When to use a cache? (hint: almost everywhere ☐)
 - Consistency and performance considerations
 - Eviction policies
 - **Systems highlight: Memcached and Redis**

Tentative Contents and Syllabus

- Tensorflow and pipelines
 - Tensorflow programming model
 - ML hardware trends
 - ML pipelines
 - Data validation and data quality
 - **Systems highlight: Tensorflow and Tensorflow Extended**
- Security and privacy
 - Security of big data systems
 - Privacy consideration
 - Data compliance and access control

Performance Concepts and Rules of Thumb



Performance Evaluation

- Metric: something we measure
- Goal: evaluate how good/bad our computer system is performing
- Examples:
 - Power consumed by our database
 - Cost of running our web application
 - Average time it takes to render a user page
 - How many users can we support at the same time
- Metrics allow us to compare two computer systems

Tradeoff: latency vs. throughput

- Pizza delivery example
 - Do you want your pizza hot?
 - Do you want your pizza to be cheap?
- Why do these conflict?
- Two different strategies for pizza company
 - Often we have a requirement for both (I want my pizza to be delivered in X time as cheaply as possible)
- Latency = execution time for a single task
- Throughput = number of tasks per unit time
- A more relevant example:
 - Assuming cars drive at 65mph, so self driving car needs to recognize an object in 0.1 seconds
 - Object recognition system needs to process 1 million object recognition tasks every second

Latency vs. Throughput is often a trade off

Plane	DC to Paris	Speed	Passengers	Throughput (pmp)
Boeing 747	6.5 hours	610 mph	470	286,700
Concorde	3 hours	1350 mph	132	178,200

■ Which plane has higher performance?

- Time to do the task (execution time)
 - **Latency**, execution time, response time
- Tasks per day, hour, week, sec (performance)
 - **Throughput**, bandwidth, operations per second

Definitions

- Performance is in units of things-per-second
 - Bigger is better
- Response time of a system Y running Z
 - $\text{performance}(Y) = \frac{1}{\text{execution time (Z on Y)}}$
- Throughput of system Y running many requests
 - $\text{performance}(Y) = \frac{\text{number of requests}}{\text{unit time}}$
- “System X is n times faster than Y” means:
 - $n = \frac{\text{performance}(X)}{\text{performance}(Y)}$

How do we improve performance?

- Suppose we have a database that processes two types of queries:
 - Query A finishes in 100 seconds
 - Query B finishes in 2 seconds
- We want better performance
 - Which query should we improve?
- The answer: it depends!

Speedup

- Make a change to the system
- Measure how much faster/slower it is

- $Speedup = \frac{Execution\ time\ before\ change}{Execution\ time\ after\ change}$

Speedup when we know details about the change

- Performance improvement depends on:

- How good is the enhancement? (factor S)
- How often is it used? (factor p)

- Speedup due to enhancement E :

- $Speedup(E) = \frac{\text{Execution time without } E}{\text{Execution time with } E} = \frac{\text{Performance with } E}{\text{Performance without } E}$
- $ExTime_{new} = ExTime_{old} * \left[(1 - p) + \frac{p}{S} \right]$
- $Speedup(E) = \frac{ExTime_{old}}{ExTime_{new}} = \frac{1}{(1-p) + \frac{p}{S}}$

Amdahl's law: example

- We built a new database that speeds up aggregate queries by 2x! Hurray!
- But.... only 10% of queries are aggregate queries

- $ExTime_{new} = ExTime_{old} * \left[(1 - p) + \frac{p}{s} \right]$

- $ExTime_{new} = ExTime_{old} * \left[(1 - p) + \frac{p}{s} \right]$

Amdahl's law in simple terms:
Make the common case fast!

- $Speedup_{total} = \frac{1}{0.95 + \frac{0.05}{2}} = 1.111$

- Amdahl's law: speedup bounded by
- Amdahl's law: speedup bounded by
- Even if aggregated queries could be completed in zero time, our **maximum** speedup would be: $\frac{1}{\text{fraction of time not enhanced}}$
- Even if aggregated queries could be completed in zero time, our **maximum** speedup would be:
- $Speedup_{optimal} = \frac{1}{0.9} = 1.111$

Useful back-of-the-envelope latency numbers (all rough estimates)

- Time measurements:
 - Nanosecond (ns): $1/1,000,000,000$ second
 - Microsecond (us): $1/1,000,000$ second
 - Millisecond (ms): $1/1000$ second
- CPU cache access: 1ns
- Memory access: 100ns
- Read a small random object from a local flash drive: 50,000ns, 50us
- Read a small object within the same region in a data center: 100,000ns, 100us
- Run a SQL query on a flash database: 1,000,000ns, 1ms
- Read a small random object from magnetic disk: 10,000,000ns, 10ms
- Run a SQL query on a disk database: 20,000,000ns, 20ms
- Roundtrip time over the internet: 100,000,000ns, 100ms
- Bounded by the speed of light!

Database example

- Scenario:
 - A user application running in the cloud needs to read a small object
 - It first checks if the object is already saved locally, either in the CPU cache or in memory:
 - 10% chance it's in the CPU cache
 - If not, 20% chance it's in memory
 - If not saved locally, it fetches it from a database from within the center network

- Compute average latency:
 - $0.1 * \text{cache latency} + 0.9 * (0.2 * \text{memory latency} + 0.8 * (\text{remote database latency}))$

 - $= 0.1\text{ns} + 18\text{ns} + 0.72 * \text{not in memory latency}$
- Remote database latency = network latency + database latency = 1,100,000ns
- Total average latency = 792,018ns or 790us
- Total average latency $\sim 0.72 * \text{not in memory latency} = 792,000\text{ns}$
- □
- **Amdahl's law: focus on the most common path taken**

Disk vs. Flash, Cost vs. Performance

- Acme runs a flu prediction service
- They have an app that displays a graph on the geographic spread of the flu, which requires running a SQL query on their database stored in the cloud
- Acme is considering running their database on flash vs. magnetic disk
 - Flash is 2X more expensive, but 20X faster
- An Acme user study shows that users don't notice page loading times, as long as they are under 300,000,000ns (300ms)
- Reminder: Internet roundtrip (100ms), disk DB access (10ms), flash DB access (1ms), read within datacenter (100us)
- Scenario 1: To compute the graph, we only need a single database access in the cloud
 - Latency with flash database: 101ms
 - **Latency with disk database: 110ms**
- Scenario 2: The app requires getting an initial response from the database, then a user input, and then another database request
 - Latency with flash database: 202ms
 - **Latency with disk database: 220ms**
- Scenario 3: The app requires 20 sequential databases accesses within the cloud to compute the graph, and then it can return
 - **Latency with flash database: 120ms**
 - Latency with disk database: 300ms

Identifying performance bottlenecks

- My application is seeing an average latency of 200ms, where is the bottleneck?
- A few guiding questions:
 1. What systems does the web page need to access? Which networks does it need to traverse?
 2. Start from the most common case + highest latency
- Example:
 - Application needs to go through the Internet $\sim 1 * 100\text{ms}$
 - Hits a server that first checks if the request is saved on memory cache in the cloud $\sim 0.2 * 100\mu\text{s}$
 - If not (80% of the time), goes over the network and accesses a single disk database $\sim 0.8 * 10\text{ms}$
- Guess 1: Internet slowdown
- Guess 2: database slowdown