

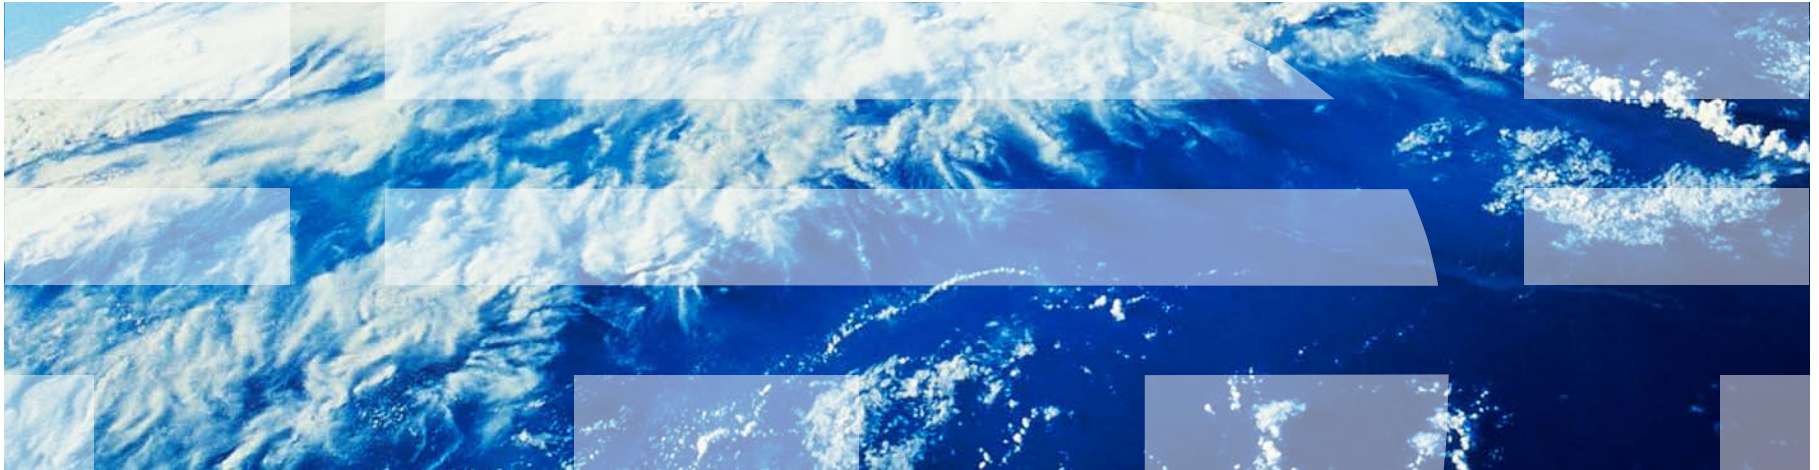


Borrowed from Shiva Shivakumar

Computer Systems for Data Science

Topic 2

Relational Model and SQL





What we'll cover in this topic

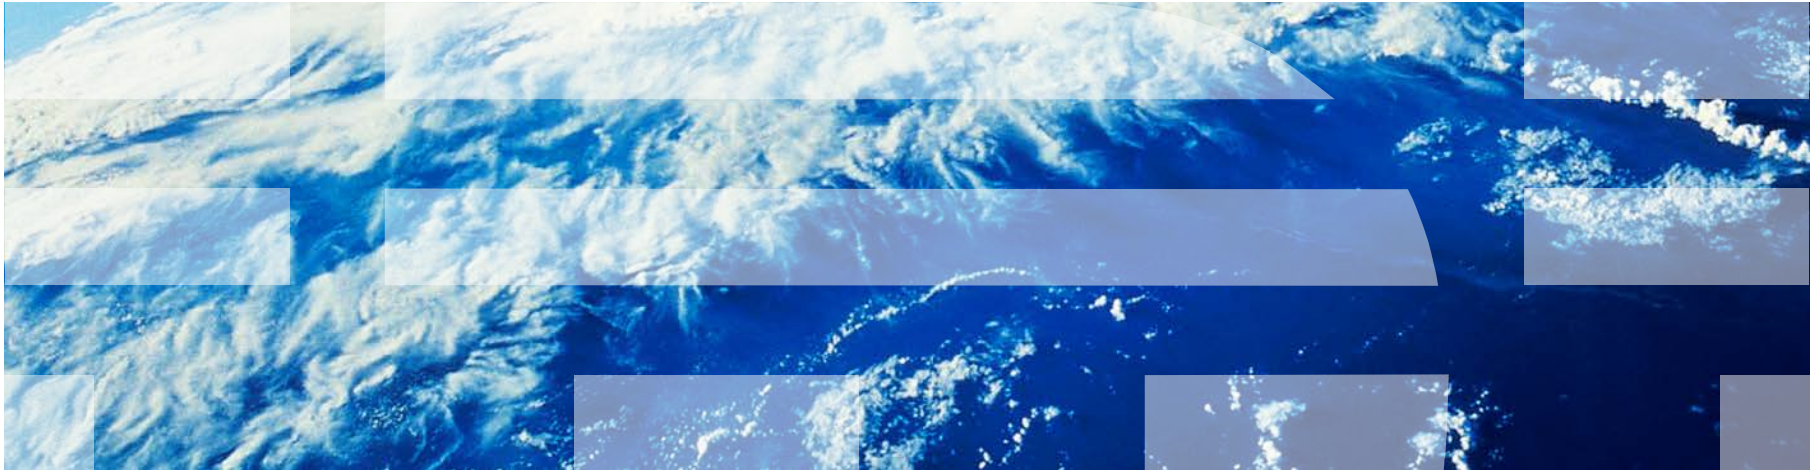
- **Intuition**

- Basic relational model
- Map-filter-reduce concept

- **Intro to SQL**

- Schemas, query structure of SELECT-FROM-WHERE, JOINS

Relational Model: Intuition

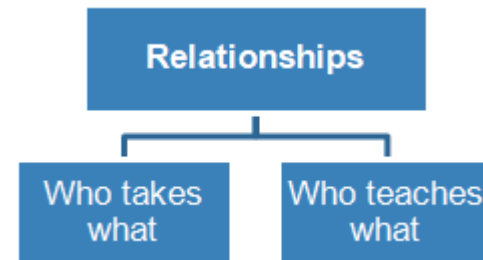
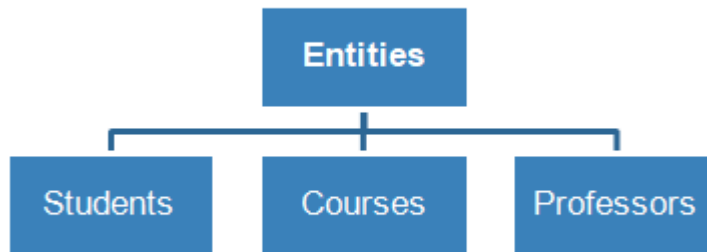


A Motivating Example

A basic Course Management System (CMS):

Entities (e.g., Students, Courses)

Relationships (e.g., Alice is enrolled in CSEE 4121)





Intuition: Spreadsheet Tables

Logical Schema

Student(cuid: string, name: string, gpa: float)

Courses(cid: string, c-name: string, room: string)

Enrolled(cuid: string, cid: string, grade: string)

Queries [“compute” over tables]

Alice’s GPA?

Jay’s classes?

AVG student GPA?

AVG student GPA in CSEE 4121?

Students			Enrolled		
CUID	Name	GPA	CUID	CID	Grade
as2121	Alice Smith	4.3	as2121	4121	A+
yg9999	Jay Goodwin	1.2	yg9999	4121	C
mc2312	Min Chang	2.2	mc2312	3292	A+
zb1111	Zorn Bjorn	3.8	zb1111	2999	D
Courses					
CID	C-Name	Room			
4121	Computer Sy	CEPSR			
3292	Databases	MUDD 1			
2999	Algorithms	MUDD 2			



Relational Model and Schemas

- Definition: Data model
 - Organizing principle of data + operations
- Relational model (aka tables)
 - Simple and most popular
 - Elegant algebra (E.F. Codd et al)
- Definition: Schema
 - Describes blueprint of table (s)
- Every relation has a schema
 - Logical Schema: describes types, names
 - Physical Schema: describes data layout
 - Virtual Schema (Views): derived tables



Data Independence

1. Can we add a new column or attribute without rewriting the application?
 - **Logical Data Independence**
 - Protection from changes in the logical structure of the data
2. Do you need to care which disks/machines are the data stored on?
 - **Physical Data Independence**
 - Protection from Physical Layout Changes



Python Operating on Lists

Basic types

- Int
- Long
- String

Map + Filter

- `map(function, list)`
- `filter(function, list)`

Map applies function to input list

Filter returns sub list that satisfies filter condition

Reduce/Aggregate

- `reduce(...)`

Reduce runs a computation on a list and returns a result
E.g., SUM, AVG, MAX



SQL Queries on Tables (Lists of Rows)

Basic types

- Int32, Int64
- Char[n]
- Float32, Float64

Map + Filter

Single Table Query

```
SELECT c1, c2
FROM T
WHERE condition;
```

Multi Table JOIN

```
SELECT c1, c2
FROM T1, T2
WHERE condition;
```

Reduce/Aggregate

```
SELECT SUM(c1*c2)
FROM T
WHERE condition
GROUP BY c3;
```

Map-Filter-Reduce pattern: Same simple/powerful idea in MapReduce, Hadoop, Spark, etc.

SQL Cheat Sheet (www.sqltutorial.org/sql-cheat-sheet)

SQL CHEAT SHEET <http://www.sqltutorial.org>



QUERYING DATA FROM A TABLE

SELECT c1, c2 FROM t;

Query data in columns c1, c2 from a table

SELECT * FROM t;

Query all rows and columns from a table

SELECT c1, c2 FROM t

WHERE condition;

Query data and filter rows with a condition

SELECT DISTINCT c1 FROM t

WHERE condition;

Query distinct rows from a table

SELECT c1, c2 FROM t

ORDER BY c1 ASC [DESC];

Sort the result set in ascending or descending order

SELECT c1, c2 FROM t

ORDER BY c1

LIMIT n OFFSET offset;

Skip *offset* of rows and return the next *n* rows

SELECT c1, aggregate(c2)

FROM t

GROUP BY c1;

Group rows using an aggregate function

SELECT c1, aggregate(c2)

FROM t

GROUP BY c1

HAVING condition;

Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

SELECT c1, c2

FROM t1

INNER JOIN t2 ON condition;

Inner join t1 and t2

SELECT c1, c2

FROM t1

LEFT JOIN t2 ON condition;

Left join t1 and t2

SELECT c1, c2

FROM t1

RIGHT JOIN t2 ON condition;

Right join t1 and t2

SELECT c1, c2

FROM t1

FULL OUTER JOIN t2 ON condition;

Perform full outer join

SELECT c1, c2

FROM t1

CROSS JOIN t2;

Produce a Cartesian product of rows in tables

SELECT c1, c2

FROM t1, t2;

Another way to perform cross join

SELECT c1, c2

FROM t1 A

INNER JOIN t2 B ON condition;

Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT c1, c2 FROM t1

UNION [ALL]

SELECT c1, c2 FROM t2;

Combine rows from two queries

SELECT c1, c2 FROM t1

INTERSECT

SELECT c1, c2 FROM t2;

Return the intersection of two queries

SELECT c1, c2 FROM t1

MINUS

SELECT c1, c2 FROM t2;

Subtract a result set from another result set

SELECT c1, c2 FROM t1

WHERE c1 [NOT] LIKE pattern;

Query rows using pattern matching %, _

SELECT c1, c2 FROM t

WHERE c1 [NOT] IN value_list;

Query rows in a list

SELECT c1, c2 FROM t

WHERE c1 BETWEEN low AND high;

Query rows between two values

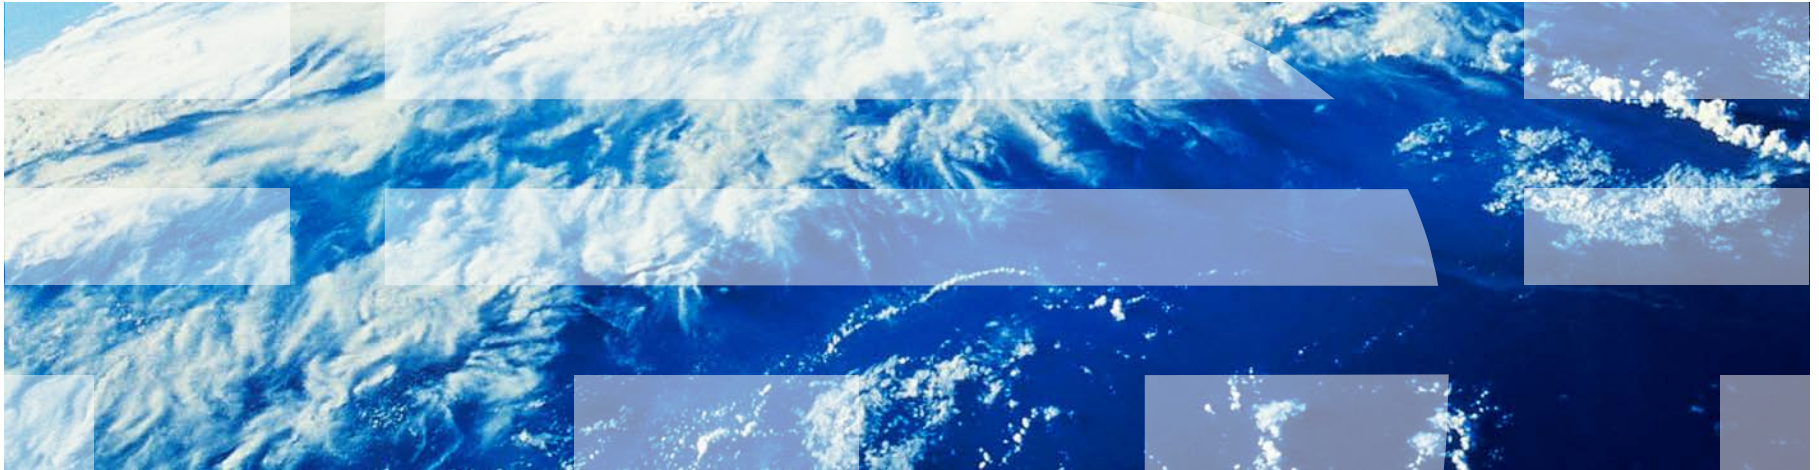
SELECT c1, c2 FROM t

WHERE c1 IS [NOT] NULL;

Check if values in a table is NULL or not



Intro to SQL





SQL Introduction

- SQL is a standard language for querying and manipulating data
- SQL is a **very high-level** programming language
 - This works because it is optimized well!

SQL stands for
Structured
Query
Language



SQL is a...

- **Data Manipulation Language (DML)**
 - Query one or more tables
 - Insert/delete/modify tuples in tables
- **Data Definition Language (DDL)**
 - Define relational schemata
 - Create/alter/delete tables and their attributes



Basic Set Algebra Concepts

- List: $[1, 1, 2, 3]$ Ordered, duplicates
- Set: $\{2, 1, 3\}$ Unordered, no duplicates
- Multiset: $\{2, 1, 3, 1\}$ Unordered, duplicates

- Unions:
 - Set: $\{2, 1, 3\} \cup \{2, 3\} = \{2, 1, 3\}$
 - Multiset: $\{2, 1, 3\} \cup \{2, 3\} = \{2, 1, 3, 2, 3\}$

- Cross-product:
 - $\{1, 1, 2, 3\} * \{y, z\} = \{1, y\}, \{1, y\}, \{2, y\}, \{3, y\}, \{1, z\}, \{1, z\}, \{2, z\}, \{3, z\}$



Tables in SQL

Product

PName	Price	Manuf
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

A relation or table is a multiset of tuples/rows having the attributes specified by the schema



Tables in SQL

Product

PName	Price	Manuf
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

An **attribute** (or **column**) is a typed data entry present in each tuple in the relation

*NB: Attributes must have an **atomic** type in standard SQL, i.e. not a list, set, etc.*

Tables in SQL

Product

PName	Price	Manuf
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

A **tuple** or **row** or **record** is a single entry in the table having the attributes specified by the schema



Tables in SQL

Product

PName	Price	Manuf
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

The number of tuples is the **cardinality** of the relation

The number of attributes is the **arity** of the relation



Data Types in SQL

Atomic types:

Characters: CHAR(20), VARCHAR(50)

Numbers: INT, BIGINT, SMALLINT, FLOAT

Others: MONEY, DATETIME...

Every attribute must have an atomic type



Table Schemas

The **schema** of a table is the table name, its attributes, and their types:

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)

A **key** is an attribute whose values are unique; we underline a key

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)



Key constraints

A **key** is a **minimal subset of attributes** that acts as a unique identifier for tuples in a relation

- A key is an implicit constraint on which tuples can be in the relation
- i.e. if two tuples agree on the values of the key, then they must be the same tuple!

Students(cuid:string, name:string, gpa: float)

1. Which would you select as a key?
2. Is a key always guaranteed to exist?
3. Can we have more than one key?



Declaring Schema

Students(cuid: *string*, name: *string*, gpa: *float*)

```
CREATE TABLE Students (  
  cuid CHAR(20),  
  name VARCHAR(50),  
  gpa float,  
  PRIMARY KEY (cuid),  
)
```



NULL and NOT NULL

- To say “don’t know the value” we use **NULL**

`Students(cuid:string, name:string, gpa: float)`

cuid	name	gpa
123	Alice	3.9
143	Jim	NULL

Say, Jim just enrolled in his first class.

In SQL, we may constrain a column to be NOT NULL, e.g., “name” in this table



General Constraints

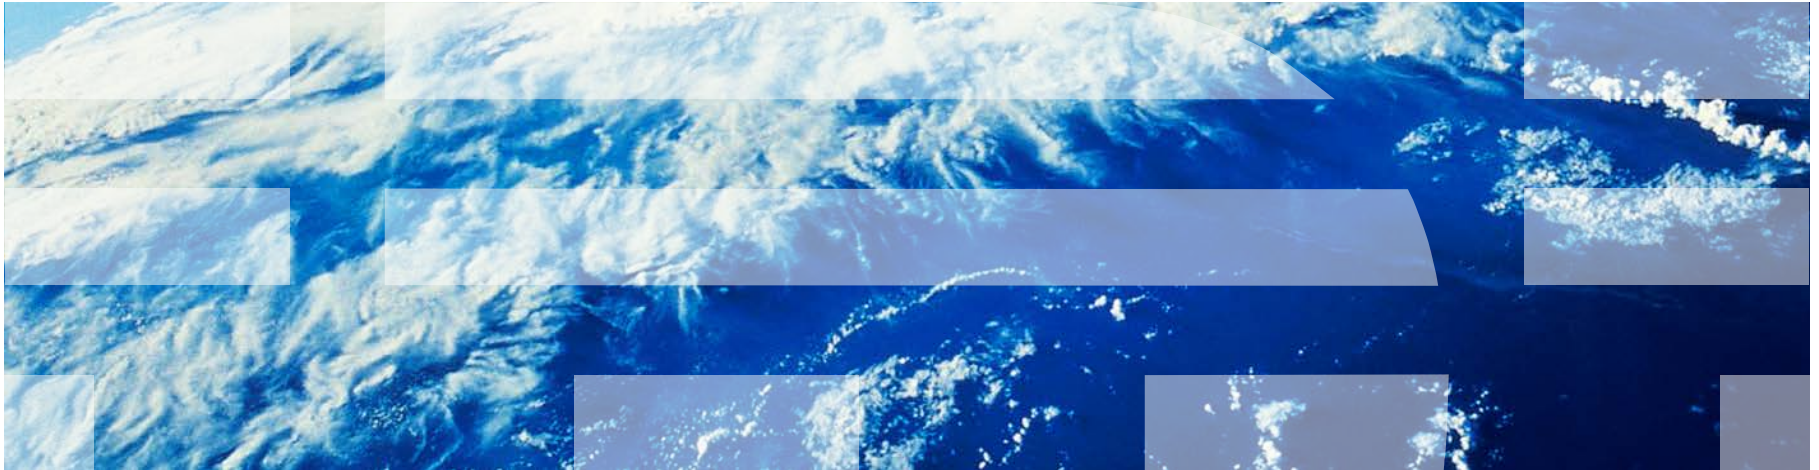
- We can actually specify arbitrary assertions
E.g. *“There cannot be 25 people in the DB class”*
- In practice, we don’t specify many such constraints. Why?
Performance!

Whenever we do something ugly (or avoid doing something convenient) it’s for the sake of performance

Summary of Schema Information

- Schema and Constraints are how databases understand the semantics (meaning) of data
- SQL supports general constraints:
Keys and foreign keys are most important

Single Table Query



SQL Query

Basic form (there are many many more bells and whistles)

SELECT <attributes>

FROM <one or more relations>

WHERE <conditions>

Call this a **SFW** query.

Simple SQL Query: Selection

Selection is the operation of filtering a relation's tuples on some condition

PName	Price	Category	Manuf
Gizmo	\$19.99	Gadgets	GWorks
Powergizmo	\$29.99	Gadgets	GWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi



```
SELECT *  
FROM Product  
WHERE Category = 'Gadgets'
```

PName	Price	Category	Manuf
Gizmo	\$19.99	Gadgets	GWorks
Powergizmo	\$29.99	Gadgets	GWorks

Simple SQL Query: Projection

Projection is the operation of producing an output table with tuples that have a subset of their prior attributes

```
SELECT Pname, Price, Manufacturer
```

```
FROM Product
```

```
WHERE Category = 'Gadgets'
```

PName	Price	Category	Manuf
Gizmo	\$19.99	Gadgets	GWorks
Powergizmo	\$29.99	Gadgets	GWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi



PName	Price	Manuf
Gizmo	\$19.99	GWorks
Powergizmo	\$29.99	GWorks

Notation

Input Schema

Product(PName, Price, Category, Manufacturer)

```
SELECT Pname, Price, Manufacturer  
FROM   Product  
WHERE  Category = 'Gadgets'
```



Output Schema

Answer(PName, Price, Manufacturer)

A Few Details

- **SQL commands** are case insensitive:
Same: SELECT, Select, select
Same: Product, product
- **Values are not:**
Different: 'Seattle', 'seattle'
- Use single quotes for constants:
'abc' - yes
"abc" - no

LIKE: Simple String Pattern Matching

```
SELECT *  
FROM Products  
WHERE PName LIKE '%gizmo%'
```

- s **LIKE** p: pattern matching on strings
- p may contain two special symbols:
 - % = any sequence of characters
 - _ = any single character

DISTINCT: Eliminating Duplicates

```
SELECT DISTINCT Category  
FROM Product
```



Category
Gadgets
Photography
Household

Versus

```
SELECT Category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

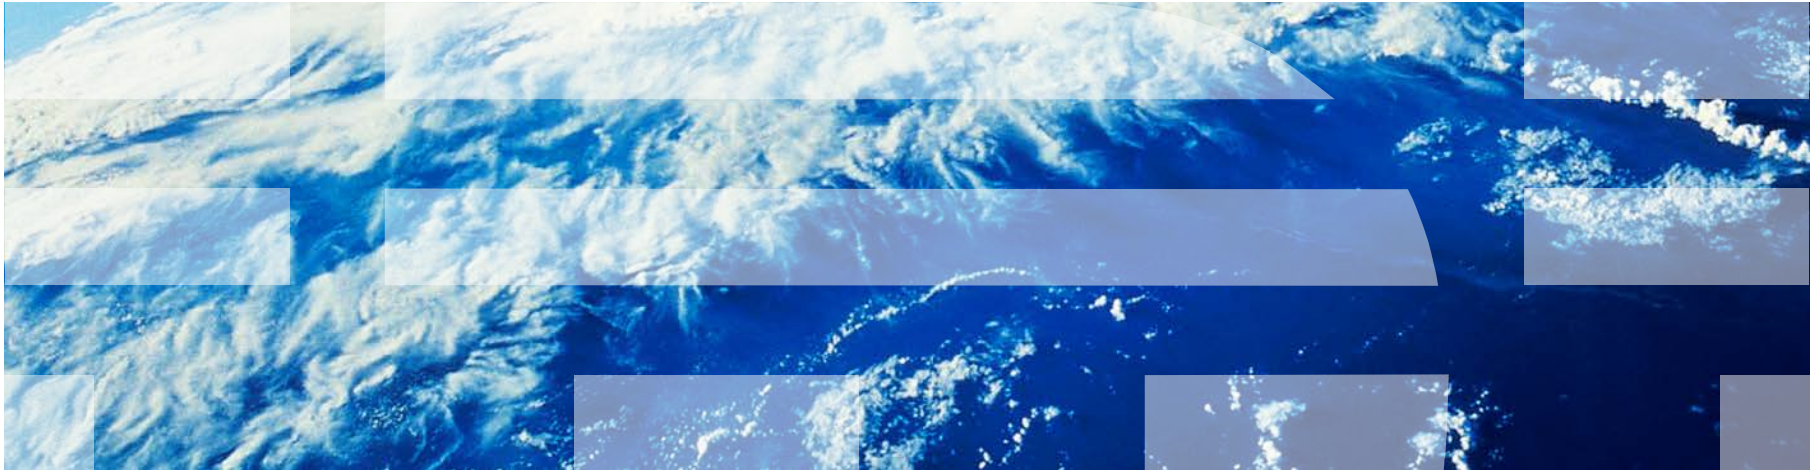
ORDER BY: Sorting the Results

```
SELECT    PName, Price, Manufacturer
FROM      Product
WHERE     Category='gizmo' AND Price > 50
ORDER BY  Price, PName
```

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.

Multi-Table Query



Foreign Key constraints

- Suppose we have the following schema :

```
Students(cuid: string, name: string, gpa: float)
Enrolled(student_id: string, cid: string, grade: string)
```

- And we want to impose the following constraint:
Only bona fide students may enroll in courses i.e. a student must appear in the Students table to enroll in a class

Students

cuid	name	gpa
102	Bob	3.9
123	Mary	3.8

Enrolled

student_id	cid	grade
102	564	A
123	537	A+

We say that cuid is a foreign key that refers to Students

Declaring Foreign Keys

Students(cuid: *string*, name: *string*, gpa: *float*)

Enrolled(student_id: *string*, cid: *string*, grade: *string*)

```
CREATE TABLE Enrolled (  
  student_id CHAR(20),  
  cid CHAR(20),  
  grade CHAR(10),  
  PRIMARY KEY (student_id, cid),  
  FOREIGN KEY (student_id) REFERENCES Students(cuid)  
)
```

Foreign Keys and update operations

Students(cuid: string, name: string, gpa: float)

Enrolled(student_id: string, cid: string, grade: string)

- What if we insert a tuple into Enrolled, but no corresponding student?

INSERT is rejected (foreign keys are constraints)!

DBA chooses

- What if we delete a student?
 1. Disallow the delete
 2. Remove all of the courses for that student
 3. SQL allows a third via NULL



Keys and Foreign Keys

Company

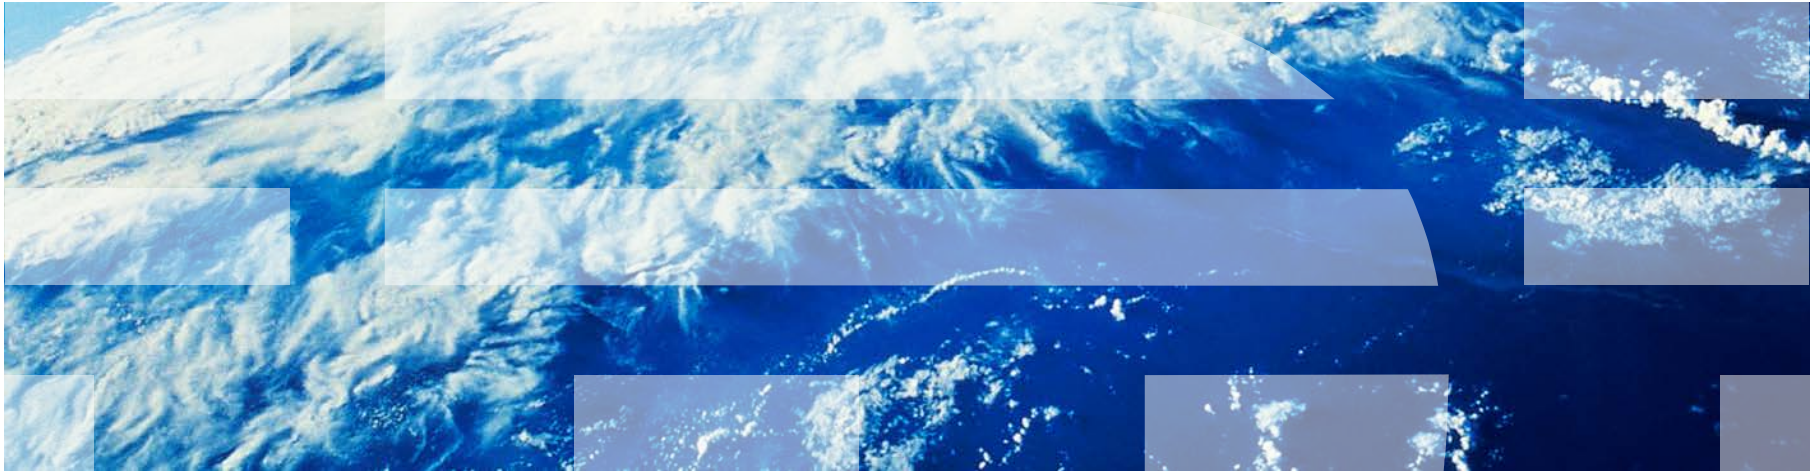
<u>CName</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

What is a foreign key vs. a key here?

Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

JOINS and Aggregations



Option 1: Organized tables, with 10s-100s of columns)

SolarPanel			
Zipcode	KW-Total	Carbon offset	...
94305	14.4	29	
94040	32.1	42	
94041	29.1	37.38	

Census		
Zipcode	Population Census	
94305	14301	
94040	20301	
94041	189	

Bike share locations			
Zipcode	Lat	Lng	
94305	35.1	122.12	
94305	35.2	122.13	
94041	35.1	121.27	
94041			

Pollution		
Zipcode	Particle count	
94305	40	
94040	22	
94041	57	



Option 2: 'FrankenTable' (with 1000s of columns)

Omnidata				SolarPanel		Census		Pollution	Bike share locations	
Zipcode	KW-Total	Carbon offset	...	Population		Particle count	Lat	Lng		
94305	14.4	29		14301		40	35.1	122.12		
94305							35.2	122.13		
94305							35.2	122.1		
94305							35.1	122.12		
94305										

????

Option 1: Organized tables, with 10s-100s of columns)

SolarPanel			
Zipcode	KW-Total	Carbon offset	...
94305	14.4	29	
94040	32.1	42	
94041	29.1	37.38	

Census		
Zipcode	Population	Census
94305		14301
94040		20301
94041		189

Bike share locations			
Zipcode	Lat	Lng	
94305	35.1	122.12	
94305	35.2	122.13	
94041	35.1	121.27	
94041			

Pollution		
Zipcode	Particle count	
94305	40	
94040	22	
94041	57	



Trade offs?

- Reads? Writes?
- 100s - thousands of applications reading/writing data

⇒ Hybrids: 1 column → all columns

(Week 7: What's a good schema design?)

Option 2: 'FrankenTable' (with 1000s of columns)

Omnidata							
SolarPanel				Census		Pollution	
Zipcode	KW-Total	Carbon offset	...	Population	Particle count	Lat	Lng
94305	14.4	29		14301	40	35.1	122.12
94305	14.4	29		14301	40	35.2	122.13
94305	14.4	29		14301	40	35.2	122.1
94305	14.4	29		14301	40	35.1	122.12

Joins

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Ex: Find all products under \$200
manufactured in Japan; return their
names and prices.

```
SELECT PName, Price
FROM   Product, Company
WHERE  Manufacturer = CName
        AND Country='Japan'
        AND Price <= 200
```

A **join**
between
tables returns
all unique
combinations
of their tuples
**which meet
some
specified join
condition**

Joins

Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19	Gadgets	GizmoWorks
Powergizmo	\$29	Gadgets	GizmoWorks
SingleTouch	\$149	Photography	Canon
MultiTouch	\$203	Household	Hitachi

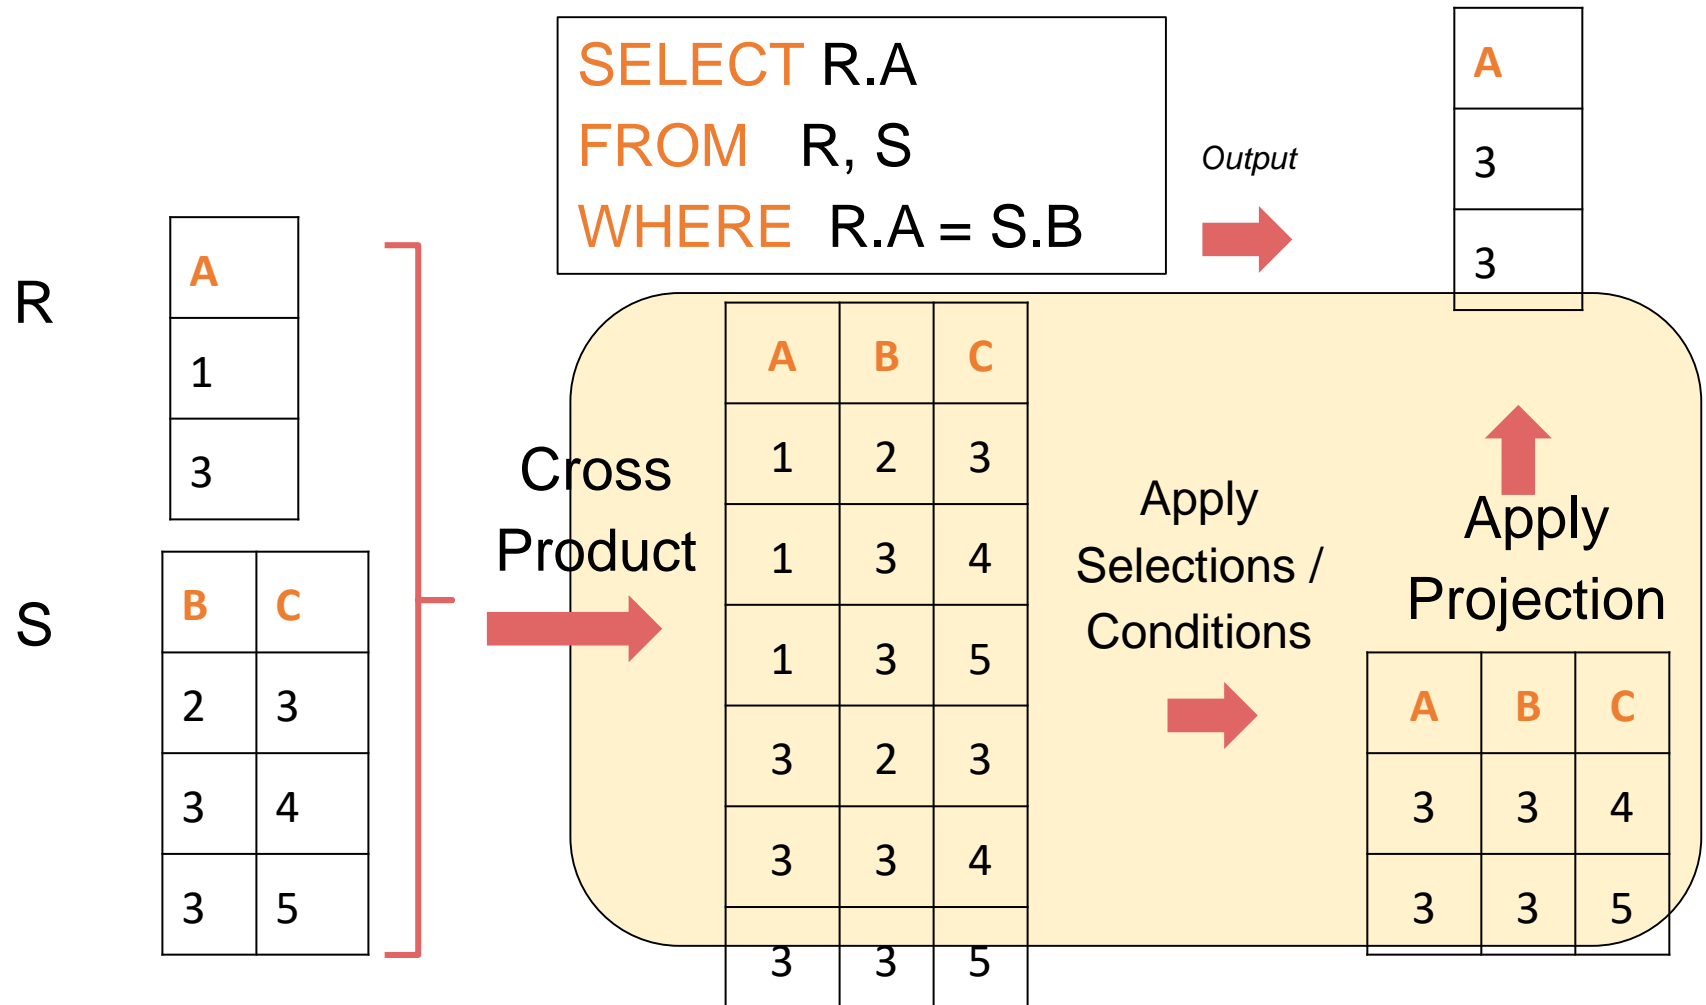
Company

<u>CName</u>	Stock Price	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

PName	Price
SingleTouch	\$149

An example of SQL semantics



Note: we say “semantics” not “execution order”

- The preceding slides show *what a join means*
- Not actually how the DBMS executes it under the covers

Aggregation

```
SELECT AVG(price)
FROM Product
WHERE maker = "Toyota"
```

```
SELECT COUNT(*)
FROM Product
WHERE year > 1995
```

- SQL supports several **aggregation** operations:
 - SUM, COUNT, MIN, MAX, AVG

*Except COUNT, all aggregations
apply to a single attribute*

More Examples

```
Purchase(product, date, price, quantity)
```

```
SELECT SUM(price * quantity)  
FROM Purchase
```

```
SELECT SUM(price * quantity)  
FROM Purchase  
WHERE product = 'bagel'
```

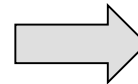
What do these mean?

Simple Aggregations

Purchase

Product	Date	Price	Quantity
bagel	10/21	1	20
banana	10/3	0.5	10
banana	10/10	1	10
bagel	10/25	1.50	20

```
SELECT SUM(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```



50 (= 1*20 + 1.50*20)

Grouping and Aggregation

Purchase(product, date, price, quantity)

```
SELECT product,  
        SUM(price * quantity) AS TotalSales  
FROM   Purchase  
WHERE  date > '10/1/2005'  
GROUP BY product
```

Find total sales
after 10/1/2005
per product.

Let's see what this means...

Grouping and Aggregation

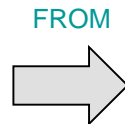
```
SELECT product,  
        SUM(price * quantity) AS TotalSales  
FROM Purchase  
WHERE date > '10/1/2005'  
GROUP BY product
```

Semantics of the query:

1. Compute the **FROM** and **WHERE** clauses
2. Group by the attributes in the **GROUP BY**
3. Compute the **SELECT** clause: grouped attributes and aggregates

1. Compute the **FROM** and **WHERE** clauses

```
SELECT product, SUM(price*quantity) AS TotalSales  
FROM Purchase  
WHERE date > '10/1/2005'  
GROUP BY product
```



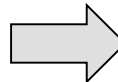
Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.50	20
Banana	10/3	0.5	10
Banana	10/10	1	10

2. Group by the attributes in the **GROUP BY**

```
SELECT product, SUM(price*quantity) AS TotalSales
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
```

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.50	20
Banana	10/3	0.5	10
Banana	10/10	1	10

GROUP BY



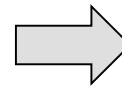
Product	Date	Price	Quantity
Bagel	10/21	1	20
	10/25	1.50	20
Banana	10/3	0.5	10
	10/10	1	10

3. Compute the **SELECT** clause: grouped attributes and aggregates

```
SELECT product, SUM(price*quantity) AS TotalSales
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
```

Product	Date	Price	Quantity
Bagel	10/21	1	20
	10/25	1.50	20
Banana	10/3	0.5	10
	10/10	1	10

SELECT



Product	TotalSales
Bagel	50
Banana	15

HAVING Clause

```
SELECT product, SUM(price*quantity)
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
HAVING SUM(quantity) > 100
```

Same query as before, except that we consider only products that have more than 100 buyers

HAVING clauses contains conditions on **aggregates**

*Whereas WHERE clauses condition on **individual tuples...***

RECAP: Joins

By default, joins in SQL are “**inner joins**”:

```
Product(name, category)
Purchase(prodName, store)
```

```
SELECT Product.name, Purchase.store
FROM Product, Purchase
WHERE Product.name = Purchase.prodName
```

```
SELECT Product.name, Purchase.store
FROM Product
JOIN Purchase ON Product.name = Purchase.prodName
```

Both equivalent:
Both INNER JOINS!

Outer Joins

- An **outer join** returns tuples from the joined relations that don't have a corresponding tuple in the other relations
 - I.e. If we join relations A and B on $a.X = b.X$, and there is an entry in A with $X=5$, but none in B with $X=5$...
 - A LEFT OUTER JOIN will return a tuple (a, NULL)!
- Left outer joins in SQL:

```
SELECT Product.name, Purchase.store  
FROM Product  
LEFT OUTER JOIN Purchase ON  
    Product.name = Purchase.prodName
```

Now we'll get products even if they didn't sell

INNER JOIN

Product

name	category
iphone	media
Tesla	car
Ford Pinto	car

Purchase

prodName	store
iPhone	Apple store
Tesla	car
iPhone	Apple store

```
SELECT Product.name, Purchase.store
FROM Product
  INNER JOIN Purchase
    ON Product.name = Purchase.prodName
```

Note: another equivalent way to write an INNER JOIN!



name	store
iPhone	Apple store
iPhone	Apple store
Tesla	car

LEFT OUTER JOIN

Product

name	category
iphone	media
Tesla	car
Ford Pinto	car

Purchase

prodName	store
iPhone	Apple store
Tesla	car
iPhone	Apple store

```
SELECT Product.name, Purchase.store
FROM Product
LEFT OUTER JOIN Purchase
ON Product.name = Purchase.prodName
```

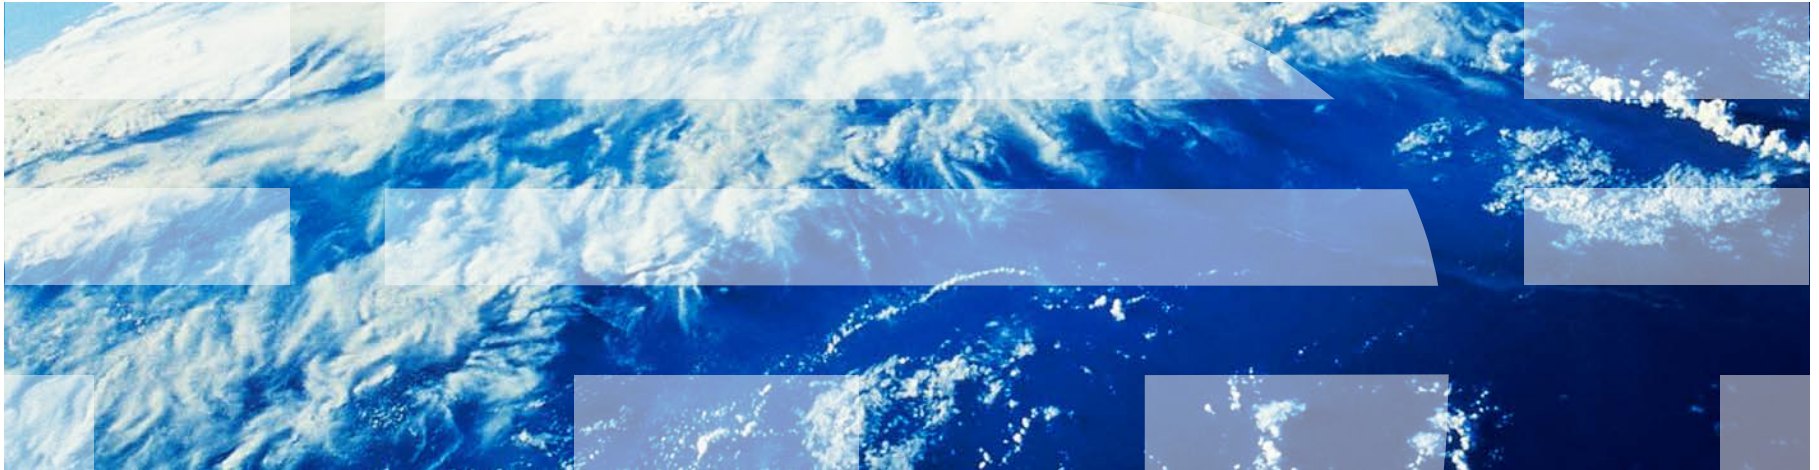


name	store
iPhone	Apple store
iPhone	Apple store
Tesla	car
Ford Pinto	NULL

Other Outer Joins

- Left outer join:
 - Include the left tuple even if there's no match
- Right outer join:
 - Include the right tuple even if there's no match
- Full outer join:
 - Include the both left and right tuples even if there's no match

Nested Queries



SQL is Compositional

Can construct powerful query chains (e.g., $f(g(\dots(x)))$)


Inputs / outputs are multisets

⇒ Output of one query can be input to another (nesting)!

⇒ Including on same table

Nested queries: Sub-queries Return Relations

Company(name, city)
Product(name, maker)
Purchase(id, product, buyer)



```
SELECT pr.maker
FROM Purchase p, Product pr
WHERE p.product = pr.name
      AND p.buyer = 'Mickey')
```

1. Companies making products bought by Mickey”
2. Location of companies?

Nested queries: Sub-queries Return Relations

Company(name, city)
Product(name, maker)
Purchase(id, product, buyer)

```
SELECT c.city
FROM   Company c
WHERE  c.name IN (
    SELECT pr.maker
    FROM   Purchase p, Product pr
    WHERE  p.product = pr.name
           AND p.buyer = 'Mickey')
```

1. Companies making products bought by Mickey”
2. Location of companies?

Subqueries Return Relations

You can also use operations of the form:

- $s > \text{ALL } R$
- $s < \text{ANY } R$
- $\text{EXISTS } R$

Ex: `Product(name, price, category, maker)`

```
SELECT name
FROM Product
WHERE price > ALL(
    SELECT price
    FROM Product
    WHERE maker = 'Gizmo-Works')
```

Find products that are more expensive than all those produced by "Gizmo-Works"

```
SELECT p1.name
FROM Product p1
WHERE p1.maker = 'Gizmo-Works'
AND EXISTS(
    SELECT p2.name
    FROM Product p2
    WHERE p2.maker <> 'Gizmo-Works'
    AND p1.name = p2.name)
```

<> means !=

Find 'copycat' products, i.e. products made by competitors with the same names as products made by "Gizmo-Works"

Note the scoping of the variables!

Example: Complex Correlated Query

```
Product(name, price, category, maker, year)
```

```
SELECT DISTINCT x.name, x.maker
FROM      Product AS x
WHERE     x.price > ALL(
    SELECT y.price
    FROM   Product AS y
    WHERE  x.maker = y.maker
           AND y.year < 1972)
```

Find products (and their manufacturers) that are more expensive than all products made by the same manufacturer before 1972

Can be very powerful (also much harder to optimize)