


# Computer Systems for Data Science


## Topic 3

### Transactions

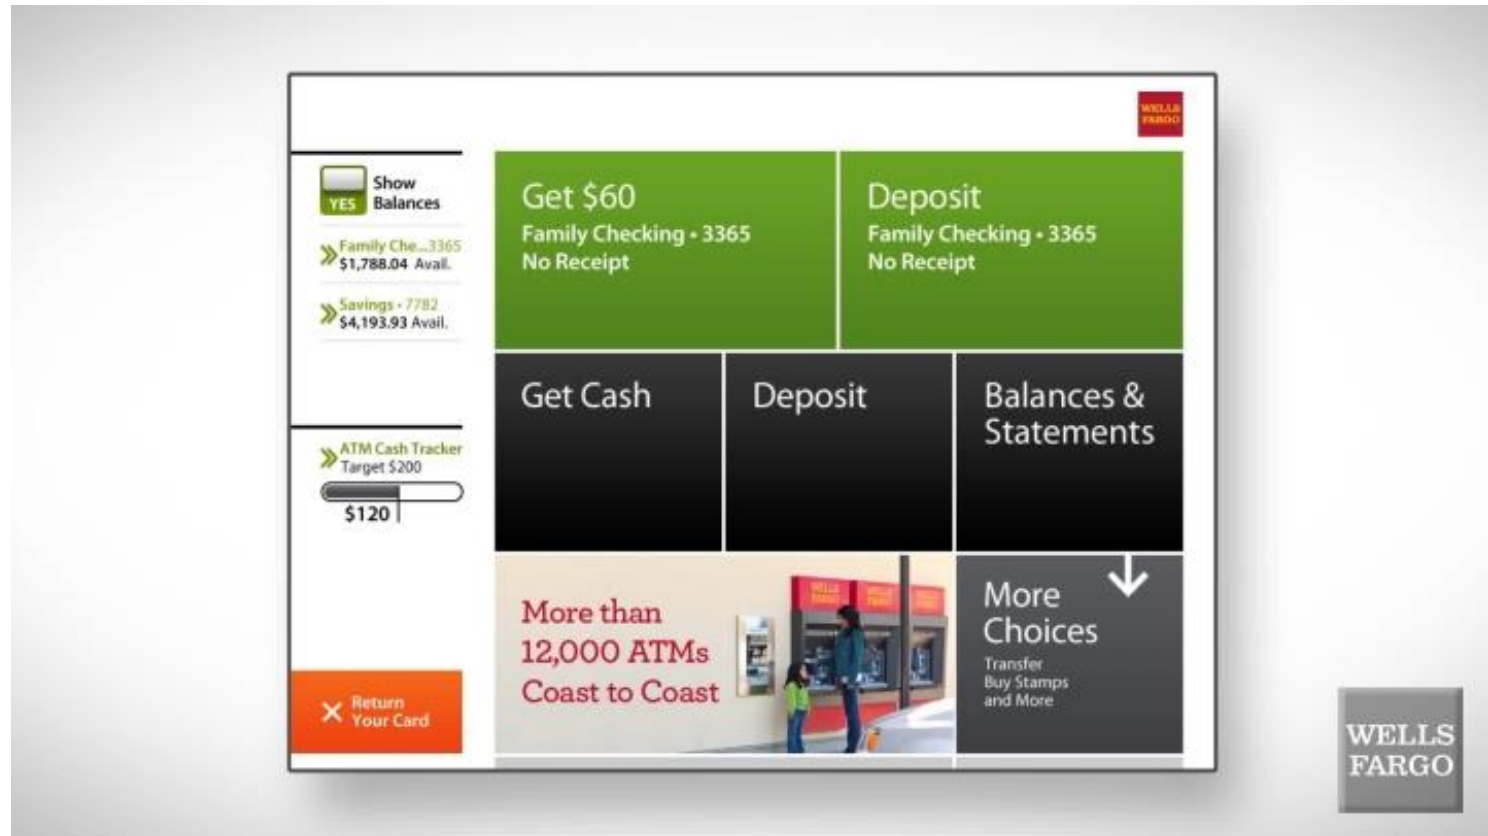


---

# Transactions



## Motivating example: an ATM



Read Balance  
Give money  
Update Balance

vs

Read Balance  
Update Balance  
Give money



It's not just about having the **correct** balance...



Visa does > 60,000 TXNs/sec with users & merchants

Want your 4\$ Starbucks transaction to wait for a 10k\$ bet in Las Vegas ?  
(Transactions can (1) be quick or take a long time, (2) unrelated to you)

# Transactions are not just used for finance




- Transactions are at the core of
- payment, stock market, banks, ticketing
  - Gmail, Google Docs (e.g., multiple people editing)



---

# Transactions



## Example: monthly bank interest transaction

Money

Account	...	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		
30108		-110
40008		110
50002		22



### 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

UPDATE Money  
SET Balance = Balance \* 1.1

## Example: monthly bank interest transaction with crash

Money

Account	...	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@10:45 am)

Account	...	Balance (\$)	
3001		550	??
4001		110	
5001		22	
6001		66	
3002		88	
4002		-200	??
5002		320	??
...			
30108		-110	
40008		110	
50002		22	??

### 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am



---

# Transactions: Basic Definition

A transaction (“TXN”) is a sequence of one or more *operations* (reads or writes) which reflects *a single real-world transition*.

**TXN either happened completely or not at all**

```
START TRANSACTION
```

```
    UPDATE Product
```

```
    SET Price = Price – 1.99
```

```
    WHERE pname = 'Gizmo'
```

```
COMMIT
```

---

# Transactions in SQL

- In “ad-hoc” SQL, each statement = one transaction
- In a program, multiple statements can be grouped together as a transaction

```
START TRANSACTION
```

```
    UPDATE Bank SET amount = amount - 100
```

```
    WHERE name = 'Bob'
```

```
    UPDATE Bank SET amount = amount + 100
```

```
    WHERE name = 'Joe'
```

```
COMMIT
```

---

# Motivation for Transactions

Grouping user actions (reads & writes) into *transactions* helps with two goals:

1. **Recovery and Durability:** Keeping the DB data consistent and durable in the face of crashes, aborts, system shutdowns, etc.
2. **Concurrency:** Achieving better performance by parallelizing TXNs *without* creating anomalies

---

# Motivation -- Recovery & Durability

1. **Recovery and durability** of user data is essential for reliable database (and other data science systems)

- The database may experience crashes (e.g. power outages, etc.)
- Individual TXNs may be aborted (e.g. by the user)

**Idea:** Make sure that TXNs are either **durably stored in full, or not at all**; keep log to be able to “roll-back” TXNs



# Protection against crashes / aborts

Client 1:

```
INSERT INTO SmallProduct(name, price)
SELECT pname, price
FROM Product
WHERE price <= 0.99
```

**Crash / abort!**

```
DELETE Product
WHERE price <=0.99
```

What goes wrong?

---

## Protection against crashes / aborts

Client 1:

```
START TRANSACTION
INSERT INTO SmallProduct(name, price)
  SELECT pname, price
  FROM Product
  WHERE price <= 0.99

DELETE Product
  WHERE price <=0.99
COMMIT OR ROLLBACK
```

Now we'd be fine! We'll see how / why this lecture

---

# Motivation -- Concurrent execution

2. **Concurrent** execution of user programs is essential for good database performance.

- Disk accesses may be frequent and slow- optimize for throughput (# of TXNs), trade for latency (time for any one TXN)
- Users should still be able to execute TXNs as if in isolation and such that consistency is maintained

**Idea:** Have the database handle running several user TXNs concurrently, in order to keep throughput high

---

## Multiple users: single statements

Client 1: **UPDATE** Product  
          **SET** Price = Price – 1.99  
          **WHERE** pname = 'Gizmo'

Client 2:   **UPDATE** Product  
          **SET** Price = Price\*0.5  
          **WHERE** pname = 'Gizmo'

Two managers attempt to discount products *concurrently*-  
What could go wrong?



---

## Multiple users: single statements

Client 1: START TRANSACTION

UPDATE Product  
SET Price = Price – 1.99  
WHERE pname = 'Gizmo'

COMMIT


Client 2: START TRANSACTION

UPDATE Product  
SET Price = Price\*0.5  
WHERE pname='Gizmo'

COMMIT

# ACID

## Atomicity, Consistency, Isolation, Durability



---

# Transaction Properties: ACID

- **A**tomic
  - State shows either all the effects of txn, or none of them
- **C**onsistent
  - Txn moves from a state where integrity holds, to another where integrity holds
- **I**solated
  - Effect of txns is the same as txns running one after another (ie looks like batch mode)
- **D**urable
  - Once a txn has committed, its effects remain in the database

ACID continues to be a source of great debate!

---

## ACID: Atomicity

- TXN's activities are atomic: all or nothing
  - Intuitively: in the real world, a transaction is something that would either occur *completely* or *not at all*
- Two possible outcomes for a TXN
  - It *commits*: all the changes are made
  - It *aborts*: no changes are made



---

## ACID: Consistency

- The tables must always satisfy user-specified *integrity constraints*
  - *Examples:*
    - Account number is unique
    - Stock amount can't be negative
    - Sum of *debits* and of *credits* is 0
- How consistency is achieved:
  - Programmer **writes** a **TXN** to go from one consistent state to a consistent state
  - *System* makes sure that the **TXN** is atomic
  - → Assuming system maintaining atomicity, this is often the user's responsibility

---

## ACID: Isolation

- A transaction executes concurrently with other transactions
- **Isolation**: the effect is as if each transaction executes in *isolation* of the others.
  - E.g. Should not be able to observe changes from other transactions during the run

---

## ACID: Durability

- The effect of a TXN must continue to exist (*“persist”*) after the TXN
  - And after the whole program has terminated
  - And even if there are power failures, crashes, etc.
  - And etc...
- Means: Write data to disk
  - And in data center settings: replicate data, backup, etc.

---

## Challenges for ACID properties

- In spite of power failures (i.e., in spite of loss of memory)
- Users may abort the program: need to “rollback changes”
  - Need to *log* what happened
- Many users executing concurrently

And all this with... Scalability and/or  
Performance!!

## A Note: ACID is contentious!


- Many debates over ACID, both **historically** and **currently**
- Some “NoSQL” DBs relax ACID
- In turn, now “NewSQL” reintroduces ACID compliance to NoSQL-style DBs...





---

# Atomicity and Durability via Logging



# Goal for this lecture: Ensuring Atomicity & Durability

## ACID

- Atomicity:
  - TXNs should either happen completely or not at all
  - If abort / crash during TXN, *no* effects should be seen

TXN 1



**Crash / abort**

**No** changes  
*persisted*

- Durability:
  - If DB stops running, changes due to completed TXNs should all persist
  - *Just store on stable disk*

TXN 2



**All** changes  
*persisted*

We'll focus on how to accomplish atomicity (via logging)

---

# Basic Idea: (Physical) Logging

## Idea:

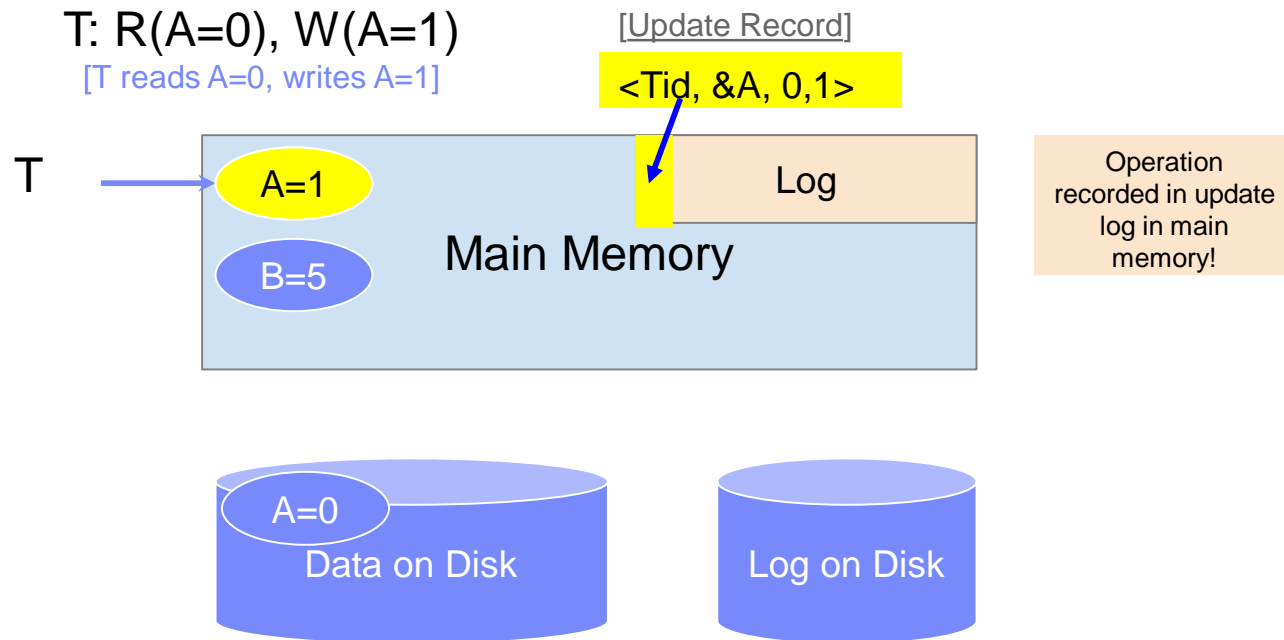
- Log consists of an ordered list of update records
- Log record contains UNDO information for every update!  
<TransactionID, location, old data, new data>

## What DB does?

- Owns the log “service” for all applications/transactions.
- Transparent to application or transaction
- Sequential writes to log, can **flush** — force writes to disk

This is sufficient to UNDO any transaction!

# A picture of logging



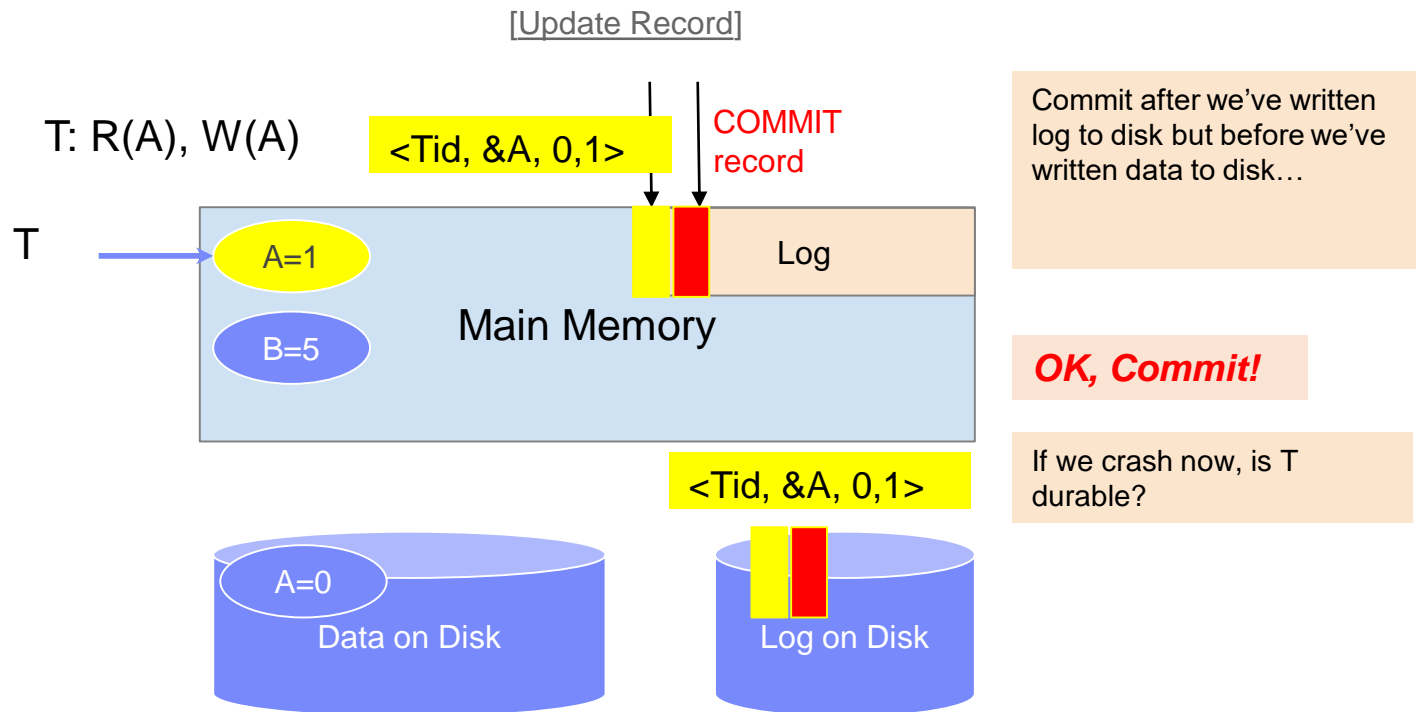
---

# Why do we need logging for atomicity?

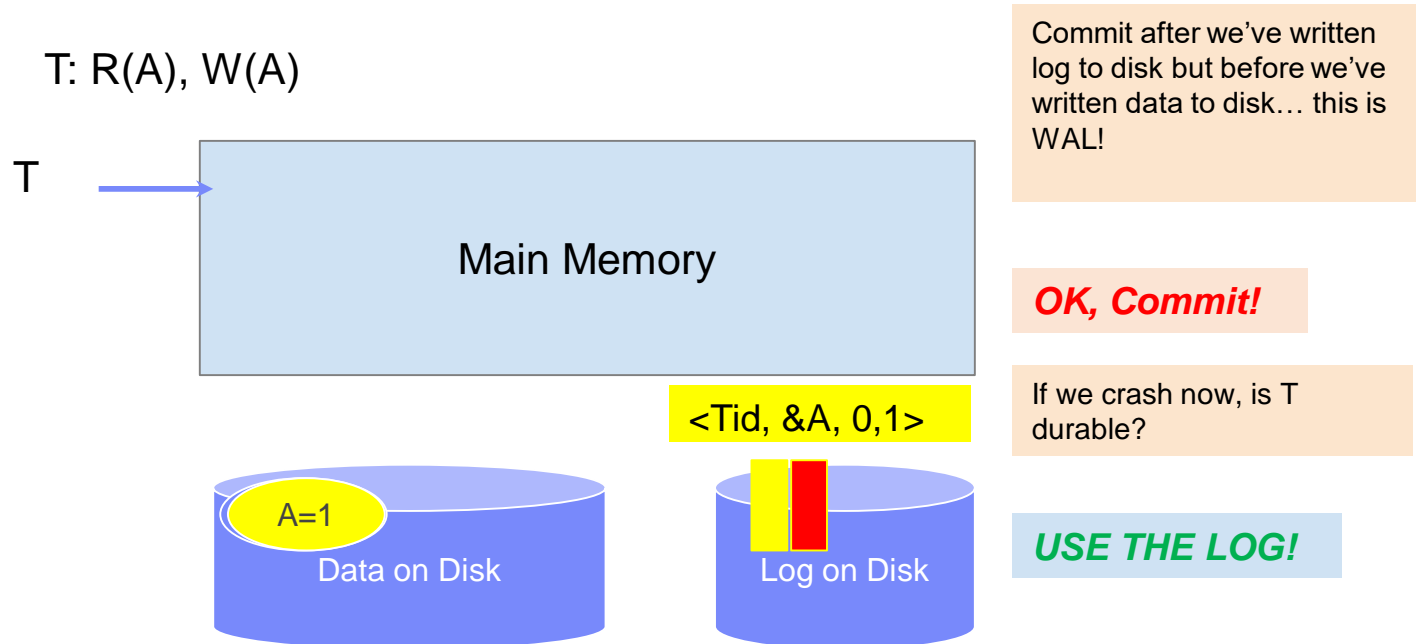
- Couldn't we just write TXN to disk **only** once whole TXN complete?
  - Then, if abort / crash and TXN not complete, it has no effect- atomicity!
  - *With unlimited memory and time, this could work...*
- However, we **need to log partial results of TXNs** because of:
  - Memory constraints (enough space for full TXN??)
  - Time constraints (what if one TXN takes very long?)

We need to write partial results to disk!  
...And so we need a **log** to be able to **undo** these partial results!

# Write-ahead Logging (WAL) Commit Protocol



# Write-ahead Logging (WAL) Commit Protocol





# Write-Ahead Logging (WAL)

## Algorithm: WAL

For each record update, write Update Record into LOG

Follow two **Flush** rules for LOG

- **Rule1:** **Flush** Update Record into LOG before corresponding data page goes to storage
- **Rule2:** Before TXN commits,
  - **Flush** all Update Records to LOG
  - **Flush** COMMIT Record to LOG

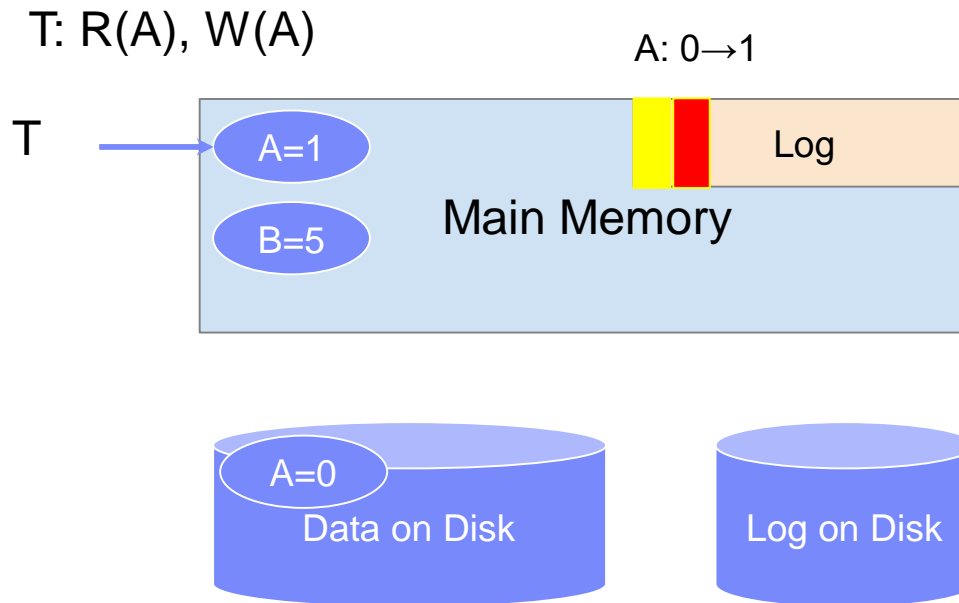
→ **Durability**

→ **Atomicity**

Transaction is committed *once COMMIT record is on stable storage*



# Incorrect Commit Protocol #1



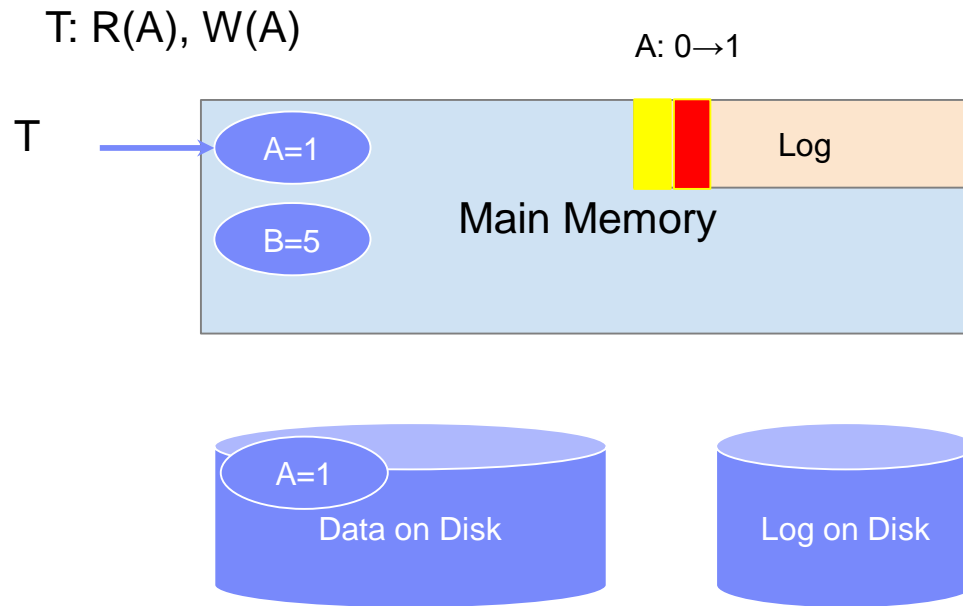
Let's try committing  
*before* we've written  
either data or log to  
disk...

**OK, Commit!**

If we crash now, is T  
durable?

**Lost T's update!**

## Incorrect Commit Protocol #2



Let's try committing *after* we've written data but *before* we've written log to disk...

**OK, Commit!**

If we crash now, is T durable? Yes! Except...

**How do we know whether T was committed??**

# Bank interest example: full run

Money			Money (@4:29 am day+1)			WAL (@4:29 am day+1)			
Account	....	Balance (\$)	Account	....	Balance (\$)				
3001		500	3001		550	T-Monthly-423	START TRANSACTION		
4001		100	4001		110	T-Monthly-423	3001	500	550
5001		20	5001		22	T-Monthly-423	4001	100	110
6001		60	6001		66	T-Monthly-423	5001	20	22
3002		80	3002		88	T-Monthly-423	6001	60	66
4002		-200	4002		-220	T-Monthly-423	3002	80	88
5002		320	5002		352	T-Monthly-423	4002	-200	-220
...		...	...		...	T-Monthly-423	5002	320	352
30108		-100	30108		-110	T-Monthly-423	...	...	...
40008		100	40008		110	T-Monthly-423	30108	-100	-110
50002		20	50002		22	T-Monthly-423	40008	100	110
						T-Monthly-423	50002	20	22
						T-Monthly-423	COMMIT		

Update  
Records

Commit  
Record

## 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

START TRANSACTION

UPDATE Money

SET Amt = Amt \* 1.10

COMMIT

# Bank interest example: with crash

Money			Money (@10:45 am)			WAL log (@10:29 am)			
Account	...	Balance (\$)	Account	...	Balance (\$)				
3001		500	3001		550	??	T-Monthly-423	START TRANSACTION	
4001		100	4001		110		T-Monthly-423	3001	500 550
5001		20	5001		22		T-Monthly-423	4001	100 110
6001		60	6001		66		T-Monthly-423	5001	20 22
3002		80	3002		88		T-Monthly-423	6001	60 66
4002		-200	4002		-200	??	T-Monthly-423	3002	80 88
5002		320	5002		320	??	T-Monthly-423	...	...
...		...	...		...		T-Monthly-423	30108	-100 -110
30108		-100	30108		-110		T-Monthly-423	40008	100 110
40008		100	40008		110		T-Monthly-423	50002	20 22
50002		20	50002		22	??	T-Monthly-423	4002	-200 -220
							T-Monthly-423	5002	320 352

## 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am

Did T-Monthly-423 complete?

Which tuples are bad?

Case1: T-Monthly-423 was crashed

Case2: T-Monthly-423 completed. 4002 deposited 20\$ at 10:45 am

# Bank interest example: with recovery

Money (@10:45 am)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		
30108		-110
40008		110
50002		22

Money (after recovery)

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

WAL log (@10:29 am)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	...	...	...
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22

System recovery (after 10:45 am)

1. Rollback uncommitted transactions
  - Restore old values from WALlog (if any)
  - Notify developers about aborted txn
1. Redo Recent transactions (w/ new values)
2. Back in business; Redo (any pending) transactions

---

## A word on performance

- Question: why is a WAL a good idea?
- Answer: updates to WAL are in sequential order!
- Recall: sequential writes are very important both for flash and magnetic disk
  - In a couple of lectures we will understand why



# An example of why sequential writes matter

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

WAL (@4:29 am day+1)

T-Monthly-423	<b>START TRANSACTION</b>		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352
T-Monthly-423	...	...	...
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	<b>COMMIT</b>		

## Cost to update all data

10M bank accounts → 10M individual random writes? (worst case)


(@ 10 ms per write for magnetic disk, that's 100,000 secs)



## Speedup for commit

100,000 secs vs 1 sec when written sequentially!!!

# Concurrency and Locking for Transactions



## Back to our bank example

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22



### Other Transactions

10:02 am Acct 3001: Wants 600\$

11:45 am Acct 5002: Wire for 1000\$

.....

.....

2:02 pm Acct 3001: Debit card for \$12.37

### 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

UPDATE Money  
SET Balance = Balance \* 1.1

Q: How do I not wait for a day to access my \$\$\$s?

---

## Big idea: locks

### ■ Intuition

- “Lock” each record for shortest time possible


### ■ Key questions

- Which records?
- For how long?
- What is the algorithm for holding them?



---

# Concurrency, Scheduling and Anomalies





---

# Concurrency: Isolation & Consistency

- DB is responsible for concurrency so that...

**Isolation** is maintained: Users must be able to execute each TXN as if they were the only user

ACID

**Consistency** is maintained: TXNs must leave the DB in consistent state

ACCID

## Example- consider two TXNs:

T1: START TRANSACTION

UPDATE Accounts

SET Amt = Amt + 100

WHERE Name = 'A'

UPDATE Accounts

SET Amt = Amt - 100

WHERE Name = 'B'

COMMIT

T1 transfers \$100 from B's account to A's account

T2: START TRANSACTION

UPDATE Accounts

SET Amt = Amt \* 1.06

COMMIT

T2 credits both accounts with a 6% interest payment

### Note:

1. DB does not care if T1  $\rightarrow$  T2 or T2  $\rightarrow$  T1 (which TXN executes first)
2. If developer does, what can they do? (Put T1 and T2 inside 1 TXN)



## Example

$T_1$

A += 100

B -= 100

T1 transfers \$100 from B's account to A's account

$T_2$

A \*= 1.06

B \*= 1.06

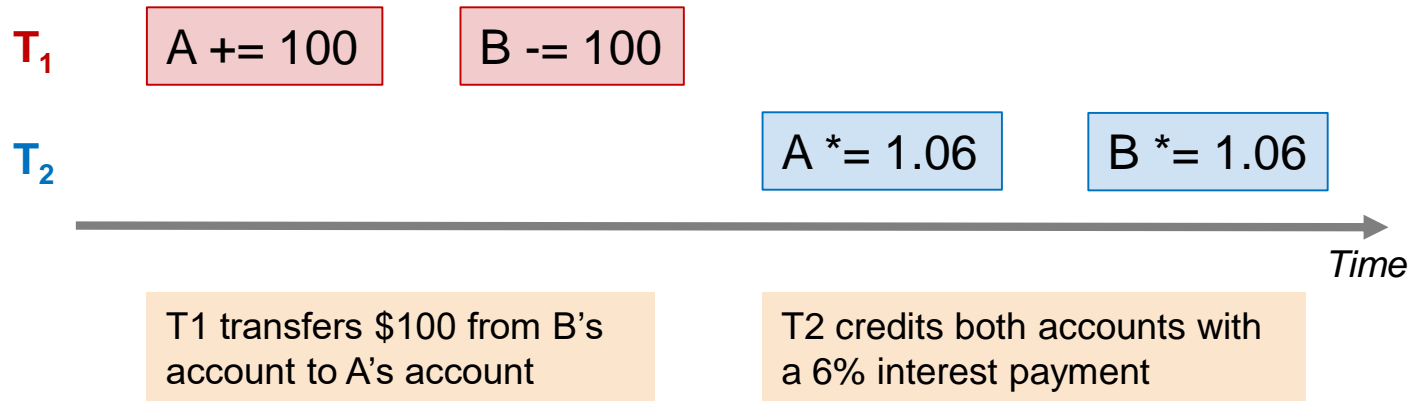
T2 credits both accounts with a 6% interest payment

Goal for scheduling transactions:

- Interleave transactions to boost performance
- Data stays in a good state after commits and/or aborts (ACID)

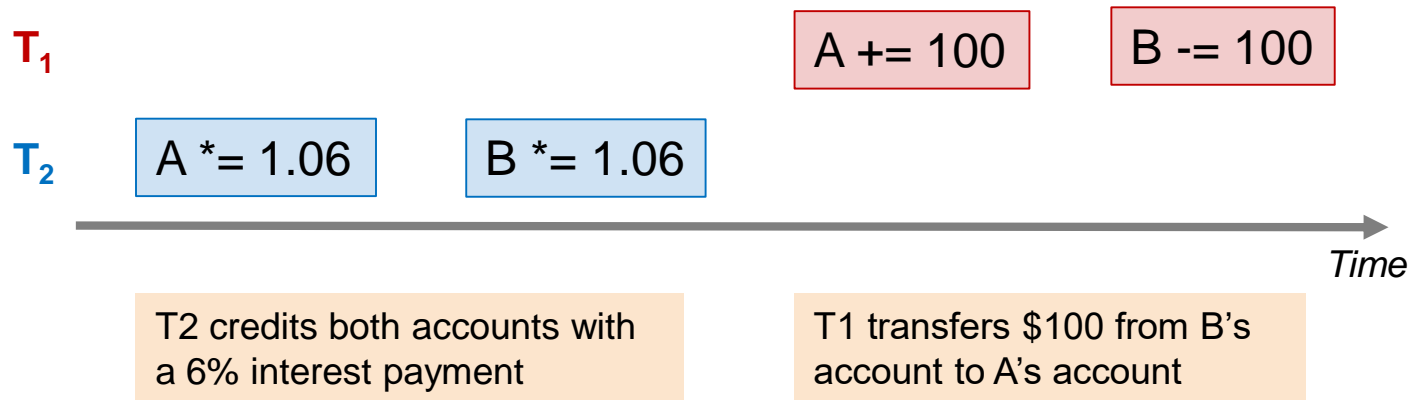
## Example- consider two TXNs:

We can look at the TXNs in a timeline view- serial execution:



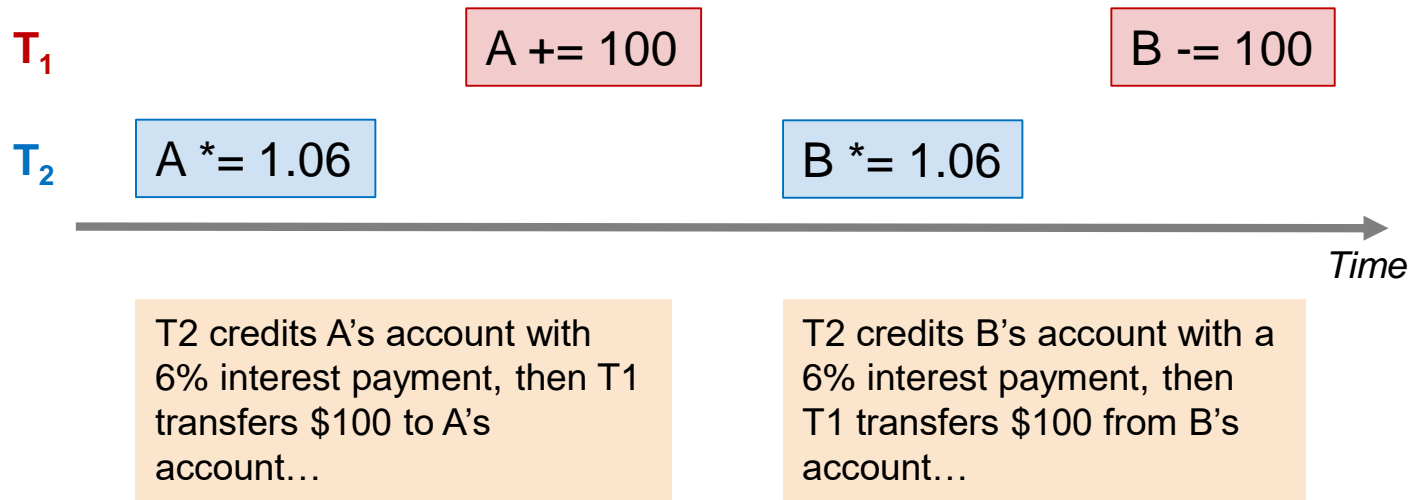
## Example- consider two TXNs:

The TXNs could occur in either order... DB allows!



## Example- consider two TXNs:

The DB can also **interleave** the TXNs



---

# Interleaving & Isolation

- The DB has freedom to interleave TXNs
- However, it must pick an interleaving or schedule such that isolation and consistency are maintained
- ⇒ Must be *as if* the TXNs had executed serially!

“With great power  
comes great  
responsibility”

**ACID**

DB must pick a schedule which maintains isolation  
& consistency

# Scheduling examples

*Starting  
Balance*

A	B
\$50	\$200

Serial schedule  $T_1, T_2$ :

**T<sub>1</sub>**      A += 100      B -= 100

$T_2$   $A^* = 1.06$   $B^* = 1.06$

A	B
\$159	\$106

**Interleaved** schedule A:

Diagram illustrating two parallel processes,  $T_1$  and  $T_2$ , each performing a sequence of operations on shared variables  $A$  and  $B$ .

- $T_1$  performs  $A += 100$  followed by  $B -= 100$ .
- $T_2$  performs  $A += 100$  followed by  $B -= 100$ .

$$T_2 \quad A^* = 1.06 \quad B^* = 1.06$$

A	B
\$159	\$106

Same result!

# Scheduling examples

*Starting  
Balance*

A	B
\$50	\$200

Serial schedule  $T_1, T_2$ :

**T<sub>1</sub>**      A += 100      B -= 100

$T_2$   $A^* = 1.06$   $B^* = 1.06$

A	B
\$159	\$106

**Interleaved** schedule B:

Diagram illustrating two parallel processes:

- T<sub>1</sub>** (Process 1): A += 100
- T<sub>2</sub>** (Process 2): B -= 100

$T_2$	$A^* = 1.06$	$B^* = 1.06$
-------	--------------	--------------

A	B
\$159	\$112

Different  
result than  
serial  
 $T_1, T_2$ !

# Scheduling examples

Serial schedule  $T_2, T_1$ :

$T_1$

A += 100

B -= 100

$T_2$

A \*= 1.06

B \*= 1.06

Starting  
Balance

A	B
\$50	\$200

A	B
\$153	\$112

**Interleaved** schedule B:

$T_1$

A += 100

B -= 100

$T_2$

A \*= 1.06

B \*= 1.06


A	B
\$159	\$112

Different  
result than  
serial  
 $T_2, T_1$   
ALSO!



# Scheduling examples

***Interleaved*** schedule B:



This schedule is different than ***any serial order!*** We say that it is **not serializable**



---

## Scheduling Definitions





- A **serial schedule** is one that does not interleave the actions of different transactions
- A and B are **equivalent schedules** if, *for any database state*, the effect on DB of executing A is **identical to** the effect of executing B
- A **serializable schedule** is a schedule that is equivalent to **some** serial execution of the transactions.

The word “**some**” makes this definition powerful & tricky!

### Serial Schedules




### Interleaved Schedules



Serial Schedules	S1, S2
Serializable Schedules	S3, S4 (And S1, S2)
Equivalent Schedules	<S1, S3> <S2, S4>
Non-serializable (Bad) Schedules	S5, S6

---


# Conflicts and Anomalies



# General DB model:


## Concurrency as Interleaving TXNs

### Serial Schedule



Each action in the TXNs  
*reads a value* and then  
*writes one back*

### Interleaved Schedule



For our purposes, having TXNs  
occur concurrently means  
**interleaving their component  
actions (R/W)**

We call the particular order  
of interleaving a **schedule**

---

# Conflict Types

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write

Thus, there are three types of conflicts:

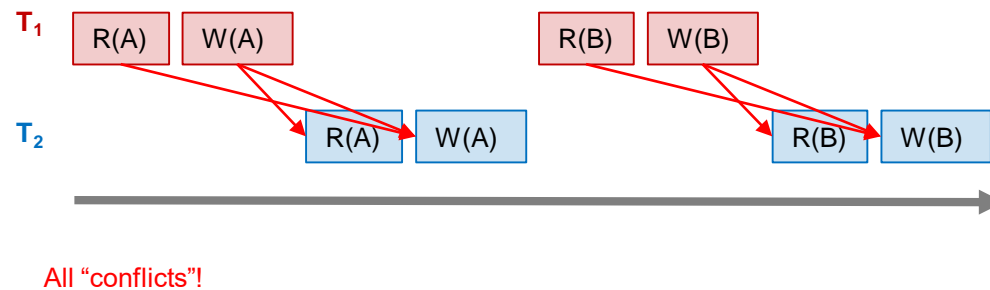
- Read-Write conflicts (RW)
- Write-Read conflicts (WR)
- Write-Write conflicts (WW)

*Why no “RR Conflict”?*

Note: **conflicts** happen often in many real world transactions. (E.g., two people trying to book an airline ticket)

# Conflicts

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write



---

## Note: Conflicts vs. Anomalies

**Conflicts** are in both “good” and “bad” schedules  
(they are a property of transactions)


Goal: Avoid Anomalies while interleaving transactions with conflicts!

- Do not create “bad” schedules where isolation and/or consistency is broken (i.e., Anomalies)



---

# Conflict Serializability



---

# Conflict Serializability

Two schedules are **conflict equivalent** if:

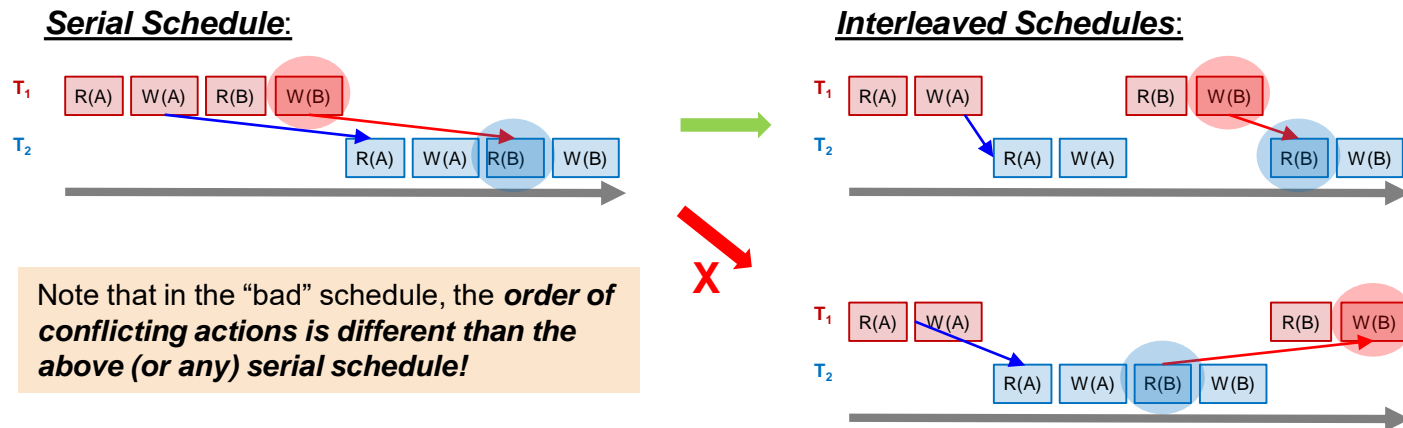
- They involve *the same actions of the same TXNs*
- Every *pair of conflicting actions* of two TXNs are *ordered in the same way*

Schedule S is **conflict serializable** if S is *conflict equivalent* to some serial schedule

**Conflict serializable  $\Rightarrow$  serializable**

So if we have conflict serializable, we have consistency & isolation!


# Example “Good” vs. “bad” schedules



Conflict serializability provides us with an operative notion of “good” vs. “bad” schedules! “Bad” schedules create data Anomalies


# The Conflict Graph

- Let's now consider looking at conflicts **at the TXN level**
- Consider a graph where the **nodes are TXNs**, and there is an edge from  $T_i \rightarrow T_j$  **if any actions in  $T_i$  precede and conflict with any actions in  $T_j$**




# What can we say about “good” vs. “bad” conflict graphs?

## Serial Schedule:




A bit complicated...

## Interleaved Schedules:




# What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:



Interleaved Schedules:



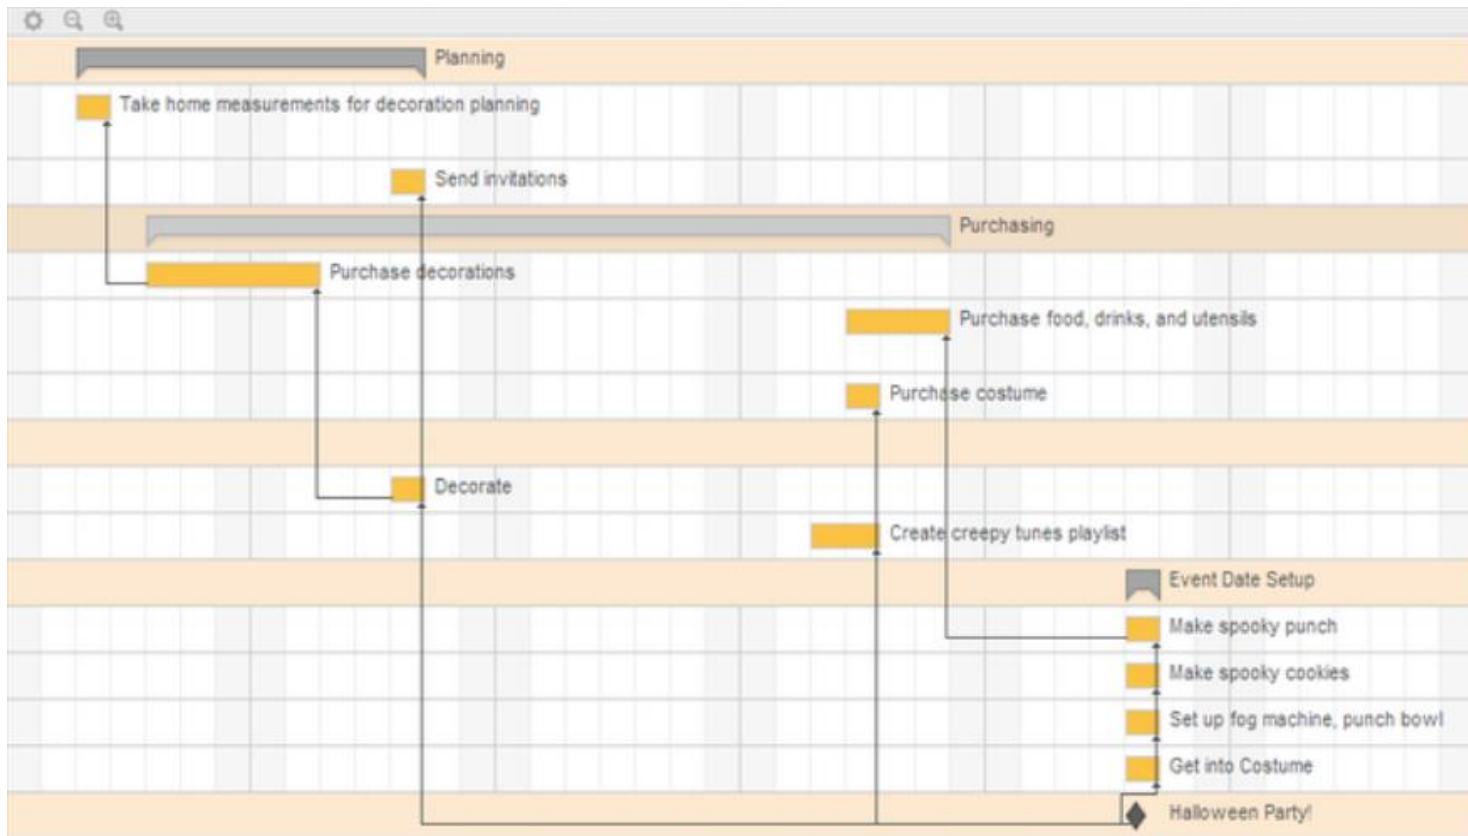
Theorem: Schedule is **conflict serializable** if and only if its conflict graph is acyclic

---

# DAGs & Topological Orderings

- A **topological ordering** of a directed graph is a linear ordering of its vertices that respects all the directed edges
- A directed **acyclic** graph (DAG) always has one or more **topological orderings**
  - (And there exists a topological ordering *if and only if* there are no directed cycles)

## Example: Project dependencies




How would you plan?  
What if there are cycles? (dependencies)



---

# DAGs & Topological Orderings

- Ex: What is one possible topological ordering here?




Ex: 0, 1, 2, 3 (or: 0, 1, 3, 2)

---

# DAGs & Topological Orderings

- Ex: What is one possible topological ordering here?



There is none!

---

## Connection to conflict serializability

- In the conflict graph, a topological ordering of nodes corresponds to **a serial ordering of TXNs**
- Thus an **acyclic** conflict graph → conflict serializable!

Theorem: Schedule is **conflict serializable** if and only if its conflict graph is **acyclic**

# Example with 5 transactions

Schedule S1


Good or Bad schedule?  
Conflict serializable?

w1(A) r2(A) w1(B) w3(C) r2(C) r4(B) w2(D) w4(E) r5(D) w5(E)

Step1  
Find conflicts  
(RW, WW, WR)

T1	w1(A)		w1(B)							
T2		r2(A)			r2(C)		w2(D)			
T3				w3(C)						
T4						r4(B)		w4(E)		
T5									r5(D)	w5(E)

Step2  
Build Conflict graph  
Acyclic?



Acyclic  
⇒ Conflict serializable!  
⇒ Serializable

Step3  
Example serial schedule  
Conflict Equiv to S1

S2

T3	T1	T4	T4	T5	T2	T2	T5	T5
w3(C)	w1(B)	r4(B)	w4(E)	r2(A)	r2(C)	w2(D)	r5(D)	w5(E)


---

## Summary

- Concurrency achieved by **interleaving TXNs** such that **isolation & consistency** are maintained
  - We formalized a notion of **serializability** that captured such a “good” interleaving schedule
- We defined **conflict serializability**

---

# 2PL: A Simple Locking Algorithm



---

## Strict Two-Phase Locking (2PL)

- Algorithm: *strict two-phase locking* - as a way to deal with concurrency
  - Guarantees conflict serializability
  - (if it completes- see upcoming...)
- Also (*conceptually*) straightforward to implement, and transparent to the user!

---

# Strict Two-phase Locking (2PL) Protocol


## TXNs obtain:

- An **X (*exclusive*) lock** on object before **writing**.
  - If a TXN holds, no other TXN can get a lock (S or X) on that object.
- An **S (*shared*) lock** on object before **reading**
  - If a TXN holds, no other TXN can get an X lock on that object
- All locks held by a TXN are released when TXN completes.

Note: Terminology here- “exclusive”, “shared”- meant to be intuitive- no tricks!



## Picture of 2-Phase Locking (2PL)



**2PL:** A transaction can not request additional locks once it releases any locks. Thus, there is a “growing phase” followed by a “shrinking phase”.

**Strict 2PL:** Release locks only at COMMIT (COMMIT Record flushed) or ABORT

---

## Strict 2PL

If a schedule follows strict 2PL, it is **conflict serializable**...

- ...and thus serializable
- ...and we get isolation & consistency!

Popular implementation

- Simple !
- Produces subset of \*all\* conflict serializable schedules
- There are MANY more complex LOCKING schemes with better performance. (See CS Database classes)
- One key, subtle problem (next)

---

# Deadlock Detection

$T_1$ 


S(A)	R(A)
------	------

$T_2$



First,  $T_1$  requests a shared lock on A to read from it

# Deadlock Detection



Next,  $T_2$  requests a shared lock on B to read from it


# Deadlock Detection: Example



$T_2$  then requests an exclusive lock on A to write to it- **now  $T_2$  is waiting on  $T_1$ ...**

Waits-For graph: Track which Transactions are waiting  
IMPORTANT: WAITS-FOR graph different than  
CONFLICT graph we learnt earlier !

# Deadlock Detection: Example



Waits-for graph:



Cycle =  
DEADLOCK

Finally,  $T_1$  requests an exclusive lock on B to write to it- **now  $T_1$  is waiting on  $T_2$ ... DEADLOCK!**

---

# Deadlocks

**Deadlock:** Cycle of transactions waiting for locks to be released by each other.

Two ways of dealing with deadlocks:

- Deadlock prevention

- Deadlock detection

---

# Deadlock Detection

Create the **waits-for graph**:

- Nodes are transactions
- There is an edge from  $T_i \rightarrow T_j$  if  $T_i$  is *waiting for  $T_j$  to release a lock*



Periodically check for (***and break***) cycles in the waits-for graph



Schedule S1

Execute with 2PL

Example with 5 Transactions (2PL)




Waits- For Graph

### Example with 5 Transactions (2PL)

Schedule S1

Execute with 2PL



Step 0

X (A)

Step 1

w1(A)

Req S(A)

Step 2

X (B)

w1(B)  
Unl B, A

Step 3

Get S(A)

Step 4

r2(A)

Step 5

S(C)

r2(C)

Step 6

S(B)

r4(B)

Step 7

X(D)

w2(D)

Unl A, C, D

Step 8

X(E)

w4(E)

Unl B,E

Step 9

S (D)

r5(D)

Step 10

X (E)

w5(E)

Unl D, E



---

# Summary

**Locking** allows only conflict serializable schedules

- If the schedule completes... (it may deadlock!)