

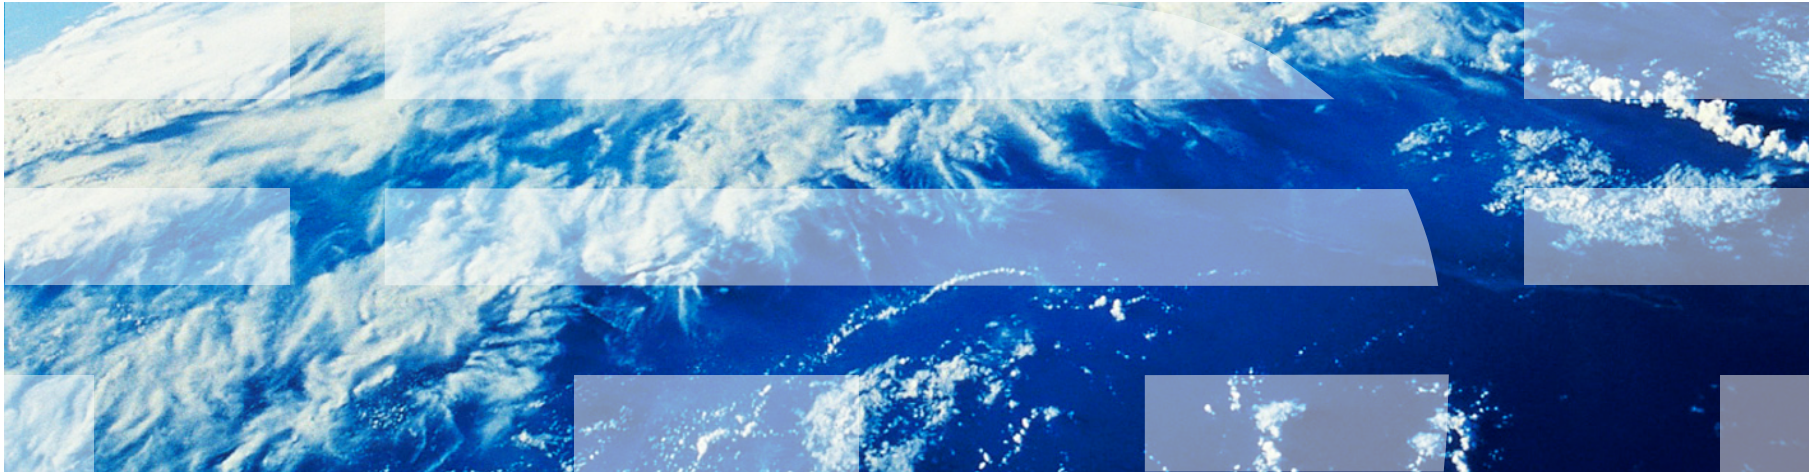
Computer Systems for Data Science

Topic 5

Distributed File Systems
Distributed Databases
Two Phase Commit



Partitioning and Replication



Introduction

- Data storage needs are growing increasingly large
 - user data at web-scale
 - 100's of millions of users, petabytes of data
 - transaction data are collected and stored for analysis.
 - multimedia objects like images/videos
- Parallel storage system requirements
 - storing large volumes of data
 - processing time-consuming decision-support queries
 - providing high throughput for transaction processing
 - Very high demands on **scalability** and **availability**

Parallel/Distributed Data Storage History

- 1980/1990s
 - Distributed database systems with tens of nodes
- 2000s:
 - Distributed file systems with 1000s of nodes
 - Millions of Large objects (100's of megabytes)
 - Web logs, images, videos, ...
 - Typically create/append only
 - Distributed data storage systems with 1000s of nodes
 - Billions to trillions of smaller (kilobyte to megabyte) objects
 - Social media posts, email, online purchases, ...
 - Inserts, updates, deletes
 - **Key-value stores**
- 2010s: Distributed database systems with 1000s of nodes

Storage Parallelism

- Reduce the time required to retrieve data from disk by partitioning the relations on *multiple disks*, on *multiple **nodes*** (computers)
 - Our description focuses on parallelism across nodes
 - Same techniques can be used across disks on a node
- **Horizontal partitioning** – tuples of a relation are divided among many nodes such that some subset of relation resides on each node.
 - Contrast with **vertical partitioning**, e.g. $r(A,B,C,D)$ with primary key A into $r1(A,B)$ and $r2(A,C,D)$

Storage Parallelism

- Partitioning techniques (number of nodes = n):

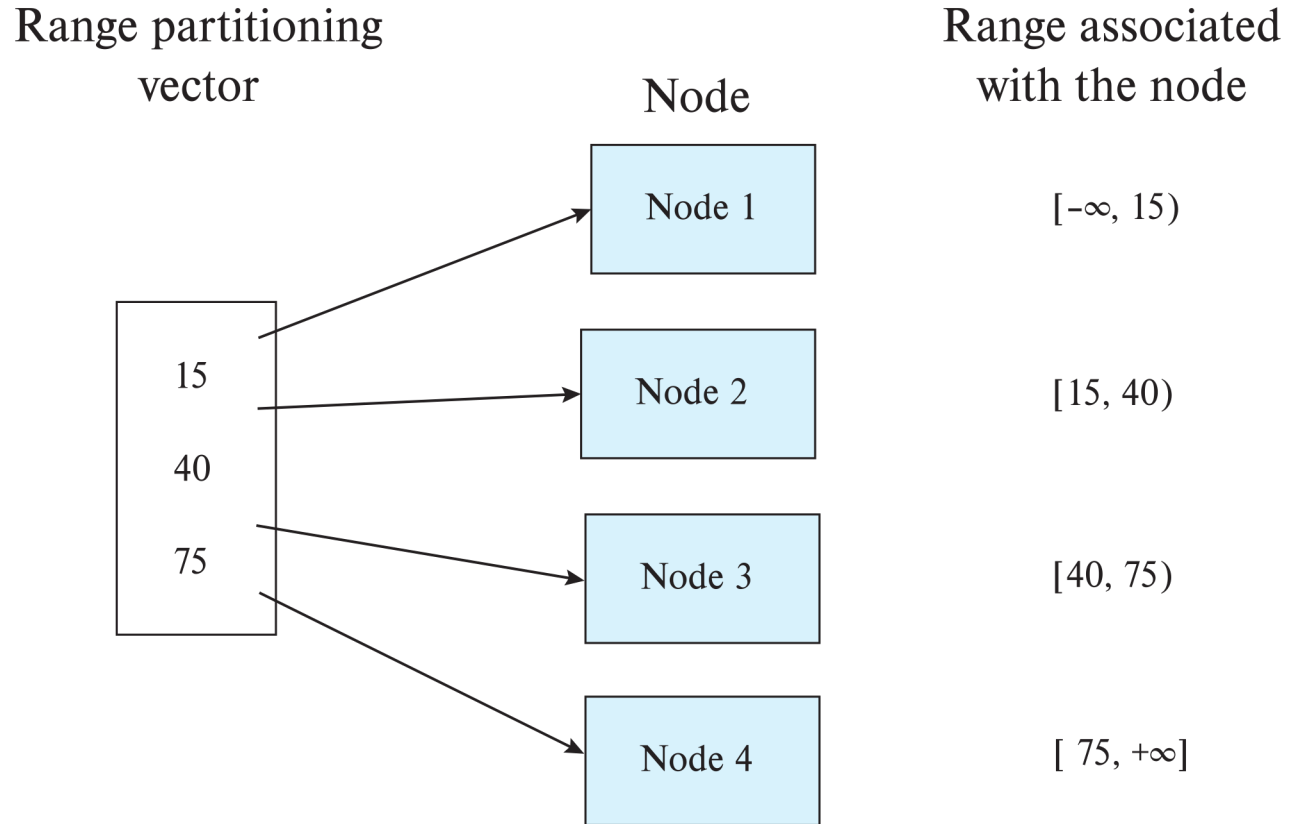
- Round-robin:**

- Send the i^{th} tuple inserted in the relation to node $i \bmod n$.

- Hash partitioning:**

- Choose one or more attributes as the partitioning attributes.
 - Choose hash function h with range $0 \dots n - 1$
 - Let i denote result of hash function h applied to the partitioning attribute value of a tuple. Send tuple to node i .

Range Partitioning



Storage Parallelism (Cont.)

Partitioning techniques (cont.):

- **Range partitioning:**

- Choose an attribute as the partitioning attribute.
- A partitioning vector $[v_0, v_1, \dots, v_{n-2}]$ is chosen.
- Let v be the partitioning attribute value of a tuple. Tuples such that $v_i \leq v_{i+1}$ go to node $i + 1$. Tuples with $v < v_0$ go to node 0 and tuples with $v \geq v_{n-2}$ go to node $n-1$.

E.g., with a partitioning vector $[5, 11]$

- a tuple with partitioning attribute value of 2 will go to node 0,
- a tuple with value 8 will go to node 1, while
- a tuple with value 20 will go to node 2.

Comparison of Partitioning Techniques

- Evaluate how well partitioning techniques support the following types of data access:
 1. Scanning the entire relation.
 2. Locating a tuple associatively – **point queries**.
 - E.g., $r.A = 25$.
 3. Locating all tuples such that the value of a given attribute lies within a specified range – **range queries**.
 - E.g., $10 \leq r.A < 25$.
- Do above evaluation for each of
 - Round robin
 - Hash partitioning
 - Range partitioning

Comparison of Partitioning Techniques

Round robin:

- Best suited for sequential scan of entire relation on each query.
 - All nodes have almost an equal number of tuples; retrieval work is thus well balanced between nodes.
- Disadvantage: all queries involving multiple tuples must be processed at all nodes

Hash partitioning:

- Similar to round robin: good for sequential access
 - Assuming hash function is good, and partitioning attributes form a key, tuples will be equally distributed between nodes
 - Less evenly distributed than round robin (uniform random is not the same as round robin)
- Disadvantage: all queries involving multiple tuples must be processed at all nodes

Comparison of Partitioning Techniques

Range partitioning:

- Provides data clustering by partitioning attribute value.
 - Good for sequential access
 - Good for point queries on partitioning attribute: only one node needs to be accessed.
- For range queries on partitioning attribute, one to a few nodes may need to be accessed
 - Remaining nodes are available for other queries.
 - Good if result tuples are from one to a few blocks.
 - But if many blocks are to be fetched, they are still fetched from one to a few nodes, and potential parallelism in disk access is wasted
 - Example of **execution skew**.

Handling Small Relations

- Partitioning not useful for small relations which fit into a single disk block or a small number of disk blocks
 - Instead, assign the relation to a single node, or
 - Replicate relation at all nodes
- For medium sized relations, choose how many nodes to partition across based on size of relation
- Large relations typically partitioned across all available nodes.

Types of Skew

- **Data-distribution skew:** some nodes have many tuples, while others may have fewer tuples. Could occur due to
 - **Attribute-value skew.**
 - Some partitioning-attribute values appear in many tuples
 - All the tuples with the same value for the partitioning attribute end up in the same partition.
 - Can occur with range-partitioning and hash-partitioning.
 - **Partition skew.**
 - Imbalance, even without attribute –value skew
 - Badly chosen range-partition vector may assign too many tuples to some partitions and too few to others.
 - Less likely with hash-partitioning

Types of Skew (Cont.)

- Note that **execution skew** can occur even without data distribution skew
 - E.g. relation range-partitioned on date, and most queries access tuples with recent dates
- Data-distribution skew can be avoided with range-partitioning by creating **balanced range-partitioning vectors**

Lecture 8



Recap of lecture 7

■ Key-value store and database example (RocksDB + MyRocks)

- Key-value stores
 - Much more limited interface than SQL (NoSQL)
 - Usually: get, put, update, delete, sometimes get_range, multiget
 - Often SQL/ACID databases are built on top of key-value stores
- RocksDB
 - Writes always go to memory: buffer writes to flash
 - Data on disk is stored in immutable files
 - Log-structure merge tree data structure compacts files, removes stale data (old/deleted key-value pairs)
 - Uses bloom filters
 - Multi-level index: outer index for each level (sparse index, in levels 1-N have distinct and sorted key-ranges), inner and dense index for each file
- MyRocks
 - Implements ACID + SQL on RocksDB
 - Uses lock table of RocksDB

Recap of lecture 7 (continued)

■ Partitioning

- Partitioning trade off
 - More nodes involved in each query: execute queries faster
 - On the hand, requires more coordination, if one node fails can slow everything down
- Round robin partitioning
 - Balances load
 - All nodes involved
- Hash partitioning
 - Usually balances load
 - All nodes involved
 - Can be used in distributed settings (without a central coordinator)
- Range partitioning
 - Can be skewed
 - Few nodes involved

■ Skew

- Data skew
 - Data is not distributed evenly (e.g., using range partitioning on attribute, and certain values are more likely than others)
- Execution skew
 - Certain data is accessed more frequently by a workload (e.g., certain rows are more “popular”/“hot”)


Today: Distributed systems

- Centralized master model
- Distributed file systems
 - Highlight: Hadoop File System (HDFS)
- Distributed transactions and databases
 - Two phase commit
 - Improvements of two phase commit
- If we have time:
- MapReduce
 - Programming model
 - Dealing with failures
- Spark intro

Logistics

- Midterm
 - You still have time to take it until tomorrow at noon
- Programming assignment #2
 - Should be available
 - Part 2 (streaming) is now optional
 - If you submit it, we will grade it, if not we won't
 - Recommended if you can — nice introduction to streaming

Master node



Routing of Queries

- Partition table typically stored at a **master** node
- Two alternative designs:
 1. Queries are sent first to **master**, which forwards them to appropriate node
 - Disadvantage: need to communicate with master for each query
 2. Master tells **client** (the node asking the query) which nodes stores which key range → clients directly communicate with data nodes
 - Advantage: do not need to talk to master for each query
 - Disadvantages: what happens if a node dies and client has old information?
- Scalability
 - Examples of systems that use a master node: Hadoop File System, Google File System, BigTable (precursor to BigQuery), HBase (open source version of BigTable)
- **Consistent hashing** is an alternative fully-distributed scheme, without a master, uses a form of hash partitioning
 - Without any master nodes, works in a completely peer-to-peer fashion
 - Advantage: no single point of failure (master), scalability
 - Disadvantage: typically requires more communication/coordination among nodes
 - Examples of systems that use consistent hashing: Cassandra, DynamoDB

Replication

- Goal: **availability** despite failures
- Data replicated at 2, often 3 nodes
 - Why 3?
- Unit of replication typically a partition (tablet)
- Requests for data at failed node automatically routed to a replica
- Partition table with each tablet replicated at two nodes

Value	Tablet ID	Node ID
2012-01-01	Tablet0	Node0,Node1
2013-01-01	Tablet1	Node0,Node2
2014-01-01	Tablet2	Node2,Node0
2015-01-01	Tablet3	Node2,Node1
2016-01-01	Tablet4	Node0,Node1
2017-01-01	Tablet5	Node1,Node0
2018-01-01	Tablet6	Node1,Node2
MaxDate	Tablet7	Node1,Node2

Basics: Data Replication

- Location of replicas
 - **Replication within a data center**
 - Handles machine failures
 - Reduces latency if copy available locally on a machine
 - Replication within/across racks
 - **Replication across data centers**
 - Handles data center failures (power, fire, earthquake, ..), and network partitioning of an entire data center
 - Provides lower latency for end users if copy is available on nearby data center

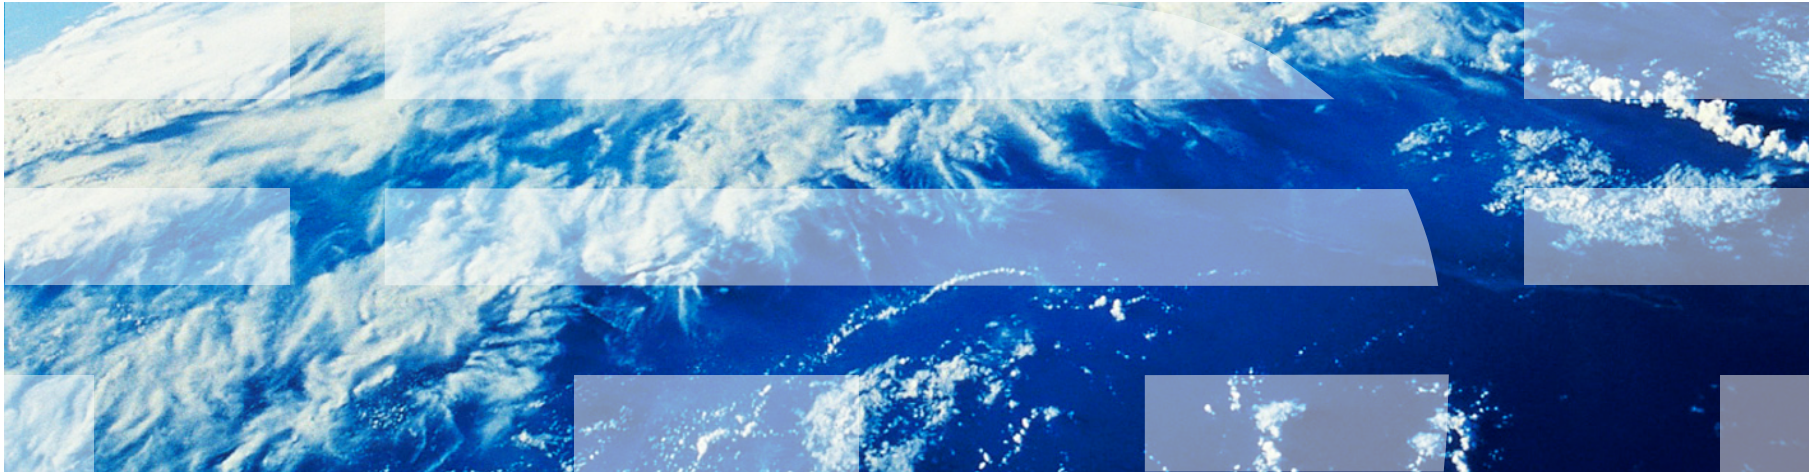
Updates and Consistency of Replicas

- Replicas must be kept consistent on update
 - Despite failures resulting in different replicas having different values (temporarily), reads must get the latest value.
 - Special concurrency control and atomic commit mechanisms to ensure consistency
- **Master replica (primary copy)** scheme
 - All updates are sent to master, and then replicated to other nodes
 - Reads are performed at master
 - But what if master fails? Who takes over? How do other nodes know who is the new master?

Protocols to Update Replicas

- *Two-phase commit*
 - Coming up!
 - Assumes all replicas are available
- *Consensus protocols*
 - Protocol followed by a set of replicas to agree on what updates to perform in what order
 - **Can work even without a designated master**

Distributed File Systems

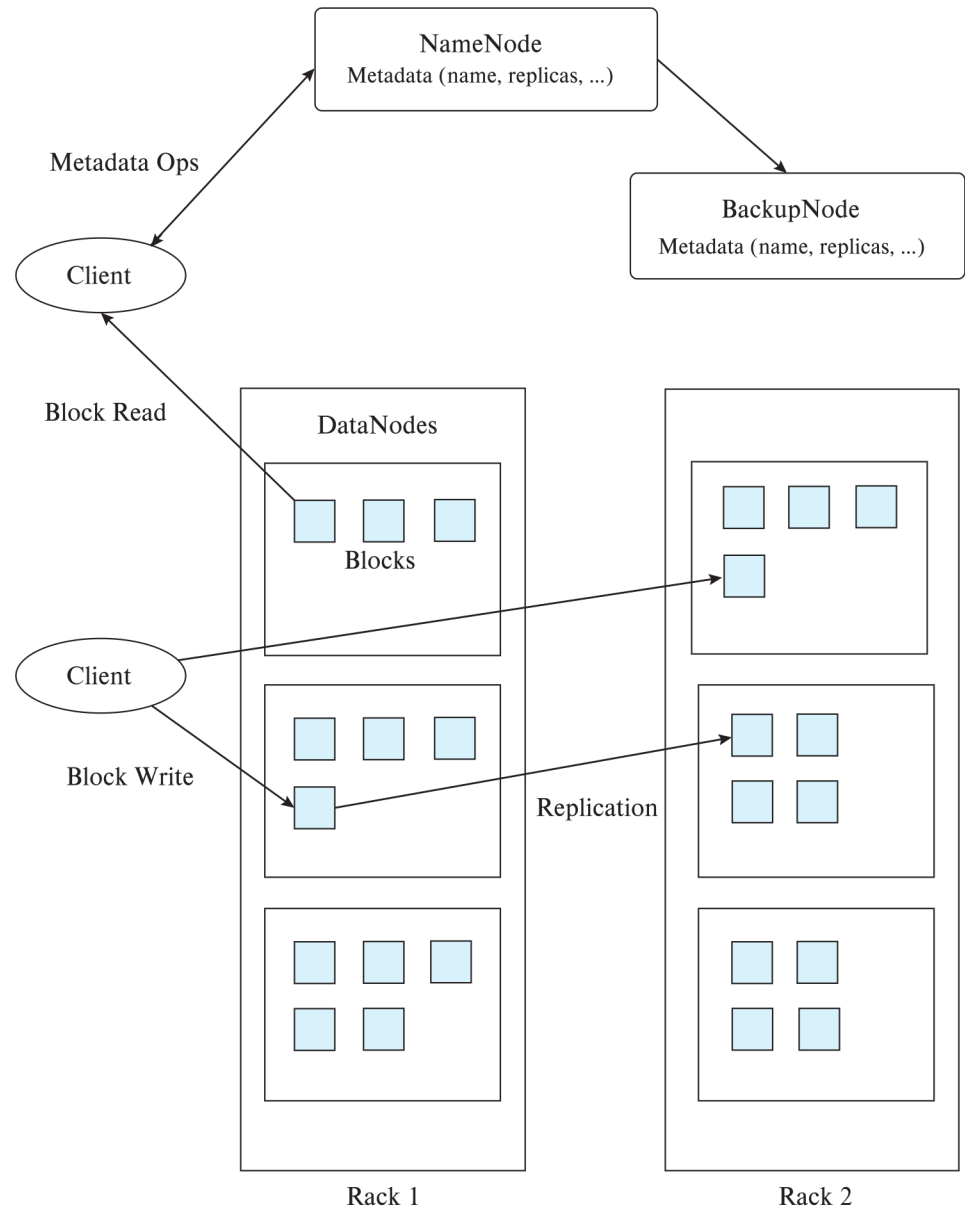


Distributed File Systems

- **Hadoop File System (HDFS)**
- Google File System (GFS)
- And older ones like CODA
- And more recent ones such as Google Colossus
- Basic architecture:
 - Master: responsible for metadata
 - Chunk servers: responsible for reading and writing large chunks of data
 - Chunks replicated on 3 machines, master responsible for managing replicas
 - Replication in GFS/HDFS is within a single data center

Hadoop File System (HDFS)

- Client: sends filename to NameNode
- NameNode (the master)
 - Maps a filename to list of Block IDs
 - Maps each Block ID to DataNodes containing a replica of the block
 - Returns list of BlockIDs along with locations of their replicas
- DataNode:
 - Maps a Block ID to a physical location on disk
 - Sends data back to client



Hadoop Distributed File System

Hadoop Distributed File System (HDFS)

- Modeled after Google File System (GFS)
- Single Namespace (e.g., single directory structure) for entire cluster
- Data model
 - Write-once-read-many access model
 - Client can only append to existing files
- Files are broken up into blocks
 - Typically 64 MB block size
 - Each block replicated on multiple (e.g., 3) DataNodes
- Client
 - Finds location of blocks from NameNode
 - Accesses data directly from DataNode
 - NameNode is not on the critical path of reads and writes

Limitations of HDFS

- Central master becomes bottleneck
 - Keep directory information in memory to avoid expensive storage reads/writes
 - Memory size limits number of files
 - What happens if it fails?
- File system directory overheads per file
 - Not appropriate for storing very large number of objects
- File systems do not provide consistency guarantees
 - File systems cache blocks locally
 - Ideal for write-once and append only data
 - Can be used as underlying storage for a data storage system
 - E.g., **BigQuery/BigTable** uses GFS/Colossus underneath, **Hbase/Spark** uses HDFS underneath

Distributed File Systems vs. Databases

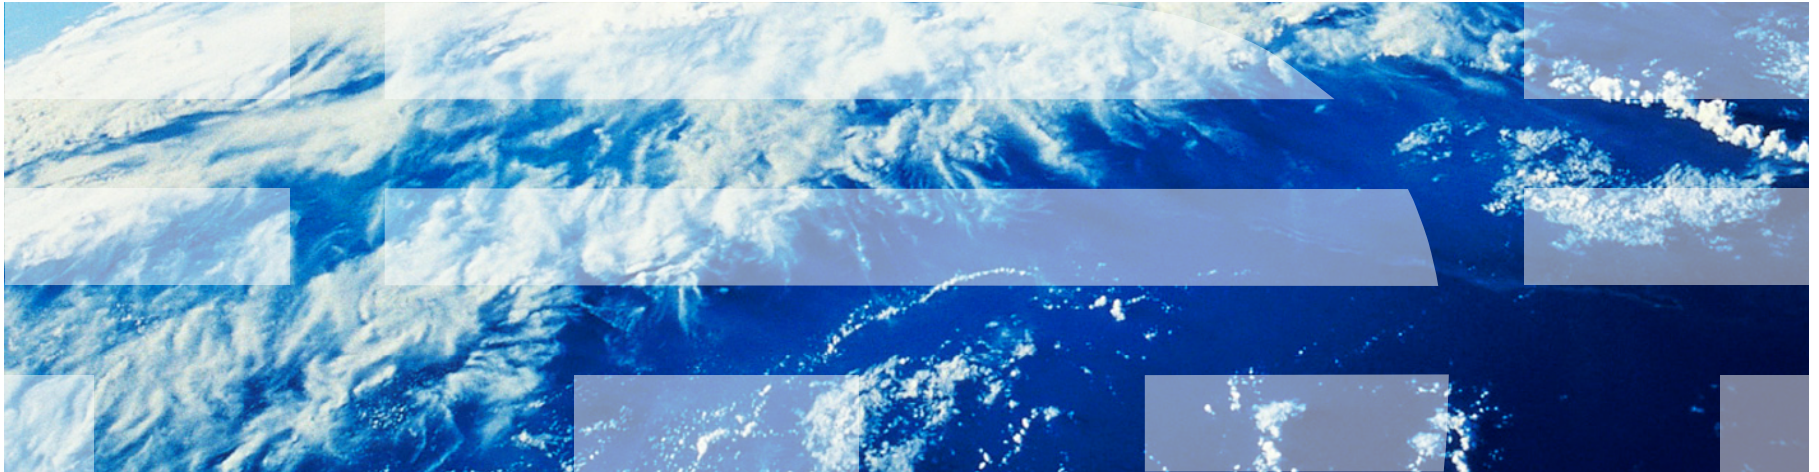
Distributed data storage implementations:

- May have limited support for relational model (no schema, or flexible schema)
- But usually do provide flexible schema and other features
 - Structured objects e.g. using JSON
 - Multiple versions of data items
- Often provide no support or limited support for transactions
 - But some do!
- Provide only lowest layer (similar to the file system layer)
- Often have KV store on top of them, and relational DB on top of the KV store

Geographically Distributed Storage

- Many storage systems today support geographical distribution of storage
 - Motivations: Fault tolerance, latency (close to user), governmental regulations
- Latency of replication across geographically distributed data centers much higher than within data center
 - Some key-value stores support **synchronous replication**
 - Must wait for replicas to be updated before committing an update
 - Others support **asynchronous replication**
 - update is committed in one data center, but sent subsequently (in a fault-tolerant way) to remote data centers
 - Must deal with small risk of data loss if data center fails.

Distributed Databases and Transactions



Approach 1: Sharding (AKA “shared-nothing” architecture)

- Divide data amongst many cheap databases (MySQL/MyRocks/PostgreSQL)
- Manage parallel access in the application
 - Partition tables map keys to nodes
 - Application decides where to route storage or lookup requests
- Scales well for both reads and writes: used widely in data center settings
- Limitations
 - Not transparent
 - application needs to be partition-aware
 - AND application needs to deal with replication
 - (Not a true parallel database, since parallel queries and transactions spanning nodes are not supported)

Approach 2: Distributed Transactions

- **Local transactions**

- Access/update data at only one database


- **Global transactions**

- Access/update data at more than one database

- Key issue: how to ensure ACID properties for transactions in a system with global transactions spanning multiple database


Distributed Transactions

- Transaction may access data at several sites.
 - Each site has a local **transaction manager**
 - Each site has a **transaction coordinator**
 - Global transactions submitted to any transaction coordinator



Distributed Transactions

- Each **transaction coordinator** is responsible for:
 - Starting the execution of transactions that originate at the site.
 - Distributing subtransactions at appropriate sites for execution.
 - Coordinating the termination of each transaction that originates at the site
 - Transaction must be committed at all sites or aborted at all sites.
- Each local **transaction manager** responsible for:
 - Maintaining a log for recovery purposes for the node it belongs to
 - Coordinating the execution and commit/abort of the transactions executing at that site.



System Failure Modes

- Failures unique to distributed systems:
 - Failure of a site
 - Loss of messages
 - Handled by networking protocols such as TCP-IP
 - **Network partition**
 - A network is said to be **partitioned** when it has been split into two or more subsystems that lack any connection between them
- Network partitioning and site failures are generally indistinguishable.

Commit Protocols

- Commit protocols are used to ensure atomicity across sites
 - A transaction which executes at multiple sites must either be committed at all the sites, or aborted at all the sites.
 - Cannot have transaction committed at one site and aborted at another
 - Why?
- **The *two-phase commit* (2PC) protocol is widely used**
- *Consensus protocols* solve a more general problem, but can be used for atomic commit
- We assume **fail-stop** model – failed sites simply stop working, and do not cause any other harm, such as sending incorrect messages to other sites (they do not become “malicious”)
 - This is the assumption in many distributed systems, but not all
 - For example, Blockchains assume that each node can be malicious (**Byzantine failure model**)
 - Generally a reasonable assumption when all the nodes reside under the same organization, with no conflicting interests

Two Phase Commit Protocol (2PC)

- Execution of the protocol is initiated by any coordinator.
- The protocol involves all the local sites at which the transaction executed
- Protocol has two phases
- Let T be a transaction initiated at site S_i , and let the transaction coordinator at S_i be C_i

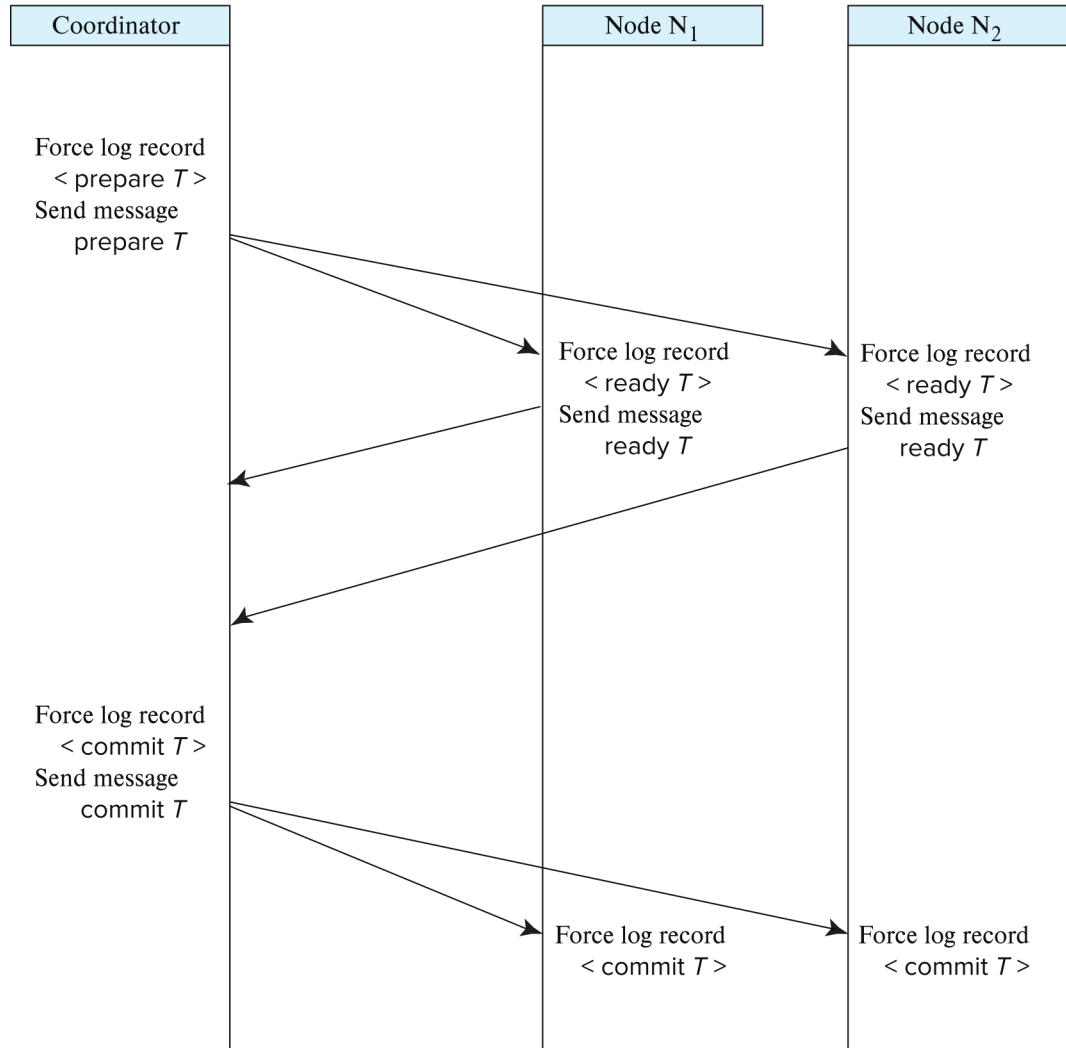
Phase 1: Obtaining a Decision

- Coordinator asks all participants to *prepare* to commit transaction T_i .
 - C_i adds the records **<prepare T >** to the log and forces log to stable storage
 - sends **prepare T** messages to all sites at which T executed
 - Upon receiving message, transaction manager at site determines if it can commit the transaction
 - if not, add a record **<no T >** to the log and send **abort T** message to C_i
 - if the transaction can be committed, then:
 - add the record **<ready T >** to the log
 - force *all records* for T to stable storage (i.e., execute the transaction locally)
 - send **ready T** message to C_i
- Transaction is now in ready state at the site

Phase 2: Recording the Decision

- T can be committed if C_i received a **ready** T message from all the participating sites: otherwise T must be aborted.
- Coordinator adds a decision record, **<commit T >** or **<abort T >**, to the log and forces record onto stable storage. Once the record is in stable storage it is irrevocable (even if failures occur)
- Coordinator sends a message to each participant informing it of the decision (commit or abort)
- Participants take appropriate action locally.
 - If the transaction was committed, they can commit locally
 - If it was aborted, they need to undo the transaction based on their local logs

Two-Phase Commit Protocol



Handling of Failures - Site Failure

When site S_k recovers, it examines its log to determine the fate of transactions active at the time of the failure.

- Log contain **<commit T >** record: site executes **redo** (T)
 - Uses the WAL
- Log contains **<abort T >** record: site executes **undo** (T)
 - Uses the WAL
- Log contains **<ready T >** record: site must consult C_i to determine the fate of T .
 - If T committed, **redo** (T)
 - If T aborted, **undo** (T)
- The log contains no control records concerning T implies that S_k failed before responding to the **prepare** T message from C_i
 - since the failure of S_k precludes the sending of such a response C_i must abort T
 - S_k must execute **undo** (T)

Handling of Failures- Coordinator Failure

- If coordinator fails while the commit protocol for T is executing then participating sites must decide on T 's fate:
 1. If an active site contains a **<commit T >** record in its log, then T must be committed.
 2. If an active site contains an **<abort T >** record in its log, then T must be aborted.
 3. If some active participating site does not contain a **<ready T >** record in its log, then the failed coordinator C_i cannot have decided to commit T . Can therefore abort T .
 4. If none of the above cases holds, then all active sites must have a **<ready T >** record in their logs, but no additional control records (such as **<abort T >** or **<commit T >**). In this case active sites must wait for C_i to recover, to find decision.
- **Blocking problem**: active sites may have to wait for failed coordinator to recover.

Handling of Failures - Network Partition

- If the coordinator and all its participants remain in one partition, the failure has no effect on the commit protocol.
- If the coordinator and its participants belong to several partitions:
 - Sites that are not in the partition containing the coordinator think the coordinator has failed, and execute the protocol to deal with failure of the coordinator.
 - No harm results, but sites may still have to wait for decision from coordinator.
- The coordinator and the sites are in the same partition as the coordinator think that the sites in the other partition have failed, and follow the usual commit protocol.
 - Again, no harm results

Recovery and Concurrency Control

- **In-doubt transactions** have a **<ready T >**, but neither a **<commit T >**, nor an **<abort T >** log record.
- The recovering site must determine the commit-abort status of such transactions by contacting other sites; this can slow and potentially block recovery.
- Recovery algorithms can note lock information in the log.
 - Instead of **<ready T >**, write out **<ready T, L >**
 - L = list of locks held by T when the log is written (read locks can be omitted).
 - For every in-doubt transaction T , all the locks noted in the **<ready T, L >** log record are reacquired.
- After lock reacquisition, transaction processing can resume; the commit or rollback of in-doubt transactions is performed concurrently with the execution of new transactions.

Avoiding Blocking During Consensus

- Blocking problem of 2PC is a serious concern
 - **Any** participant that fails or is delayed will block all other nodes!
- Idea: involve multiple nodes in decision process, so failure of a few nodes does not cause blocking as long as majority don't fail
- More general form: **distributed consensus problem**
 - A set of n nodes need to agree on a decision
 - Inputs to make the decision are provided to all the nodes, and then each node votes on the decision
 - The decision should be made in such a way that all nodes will “learn” the same value for the even if some nodes fail during the execution of the protocol, or there are network partitions.
 - Further, the distributed consensus protocol should not block, as long as a **majority** of the nodes participating remain alive and can communicate with each other
- Several consensus protocols, Paxos and Raft are popular

Using Consensus to Avoid Blocking

- After getting response from 2PC participants, coordinator can initiate distributed consensus protocol by sending its decision to a set of participants who then use consensus protocol to commit the decision
 - If coordinator fails before informing all consensus participants
 - Choose a new coordinator, which follows 2PC protocol for failed coordinator
 - If a commit/abort decision was made as long as a majority of consensus participants are accessible, decision can be found without blocking
 - If consensus process fails (e.g., split vote), restart the consensus
 - Split vote can happen if a coordinator send decision to some participants and then fails, and new coordinator send a different decision
- The **three phase commit** protocol is an extension of 3PC which avoids blocking under certain assumptions
 - Ideas are similar to distributed consensus.
- Consensus is also used to ensure consistency of replicas of a data item
 - For example, can be used to protect against failure of master nodes (e.g., the NameNode in HDFS)